

Stoneburner, Kurt

• DSC 650 - Week 06

• Assignment 6.1 - ConvNet Model that classifies images in the MNIST digital dataset.

This work is essentially copy/pasted from the book's code. I spent a fair bit of time getting my head wrapped around conv2d and Max Pooling. I added comments to clarify my current understanding and help with future reference.

```
In [1]: 1 #!/*** If this code works, then Directml and GPU operations are working
2 import tensorflow.compat.v1 as tf
3
4 tf.enable_eager_execution(tf.ConfigProto(log_device_placement=True))
5
```

```
Executing op Add in device /job:localhost/replica:0/task:0/device:DML:
0
tf.Tensor([4. 6.], shape=(2,), dtype=float32)
```

Reference: <https://keras.io/api/datasets/mnist/> (<https://keras.io/api/datasets/mnist/>)

Max-Pooling Explained: <https://analyticsindiamag.com/max-pooling-in-convolutional-neural-network-and-its-features/> (<https://analyticsindiamag.com/max-pooling-in-convolutional-neural-network-and-its-features/>)

Conv2D Official Documentation: https://keras.io/api/layers/convolution_layers/convolution2d/ (https://keras.io/api/layers/convolution_layers/convolution2d/)

```
In [8]: 1 import os
2 import sys
3 #!/*** Imports and Load Data
4 import matplotlib.pyplot as plt
5 import numpy as np
6 import pandas as pd
7
8 #!/*** Use the whole window in the IPYNB editor
9 from IPython.core.display import display, HTML
10 display(HTML("<style>.container { width:100% !important; }</style>"))
11
12 #!/*** Maximize columns and rows displayed by pandas
13 pd.set_option('display.max_rows', 100)
14 pd.set_option('display.max_columns', None)
15
16 import tensorflow as tf
17
18 from tensorflow import keras
19 from tensorflow.keras import layers
```

Load the MNIST dataset.

This is a dataset of 60,000 28x28 grayscale images of the 10 digits, along with a test set of 10,000 images.

Returns

Tuple of NumPy arrays: (x_train, y_train), (x_test, y_test).

x_train: uint8 NumPy array of grayscale image data with shapes (60000, 28, 28), containing the training data. Pixel values range from 0 to 255.

y_train: uint8 NumPy array of digit labels (integers in range 0-9) with shape (60000,) for the training data.

x_test: uint8 NumPy array of grayscale image data with shapes (10000, 28, 28), containing the test data. Pixel values range from 0 to 255.

y_test: uint8 NumPy array of digit labels (integers in range 0-9) with shape (10000,) for the test data.

```
In [10]: 1
          2 (x_train, y_train), (x_test, y_test) = tf.keras.datasets.mnist.load_d
          3 assert x_train.shape == (60000, 28, 28)
          4 assert x_test.shape == (10000, 28, 28)
          5 assert y_train.shape == (60000,)
          6 assert y_test.shape == (10000,)
          7
```

```
In [11]: 1
          2
          3 #!/*** Inputs reflects the shape of each individual piece of data.
          4 #!/*** The MNIST is a 28x28 single channel image.
          5 #!/*** The third is 1 channel. This is a greyscale image, therefore i
          6 #!/*** See Link above for further explanation
          7 inputs = keras.Input(shape=(28, 28, 1))
          8
          9 #!/*** Conv2D: Filters defines the number of tensors at the layer. Ke
         10 #!/*** Conv2D reduces the image size by (filter_size - 1) in
         11 #!/*** MaxPooling2D: Is a form of feature reduction. In this case, Th
         12 #!/*** pool-size value (2x2) is kept. With a pool value
         13 #!/*** by 75%.
         14 #!/*** At each stage of max pooling, the image gets sma
         15 #!/*** for relationships between the reduced features.
         16
         17 x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(input
         18 x = layers.MaxPooling2D(pool_size=2)(x)
         19 x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
         20 x = layers.MaxPooling2D(pool_size=2)(x)
         21 x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
         22 x = layers.Flatten()(x)
         23 outputs = layers.Dense(10, activation="softmax")(x)
         24 model = keras.Model(inputs=inputs, outputs=outputs)
         25
```

```
26 model.summary()
Model: "model_1"
```

Layer (type)	Output Shape	Param #
=====		
input_2 (InputLayer)	[(None, 28, 28, 1)]	0
conv2d_3 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_2 (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_4 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_3 (MaxPooling2D)	(None, 5, 5, 64)	0
conv2d_5 (Conv2D)	(None, 3, 3, 128)	73856
flatten_1 (Flatten)	(None, 1152)	0
dense_1 (Dense)	(None, 10)	11530
=====		
Total params: 104,202		
Trainable params: 104,202		
Non-trainable params: 0		
=====		

```
In [14]: 1 x_train = x_train.reshape((60000, 28, 28, 1))
          2 x_train = x_train.astype("float32") / 255
          3 x_test = x_test.reshape((10000, 28, 28, 1))
          4 x_test = x_test.astype("float32") / 255
          5 model.compile(optimizer="rmsprop",
          6                 loss="sparse_categorical_crossentropy",
          7                 metrics=["accuracy"])
          8 model.fit(x_train, y_train, epochs=5, batch_size=64)
```

```
Epoch 1/5
938/938 [=====] - 21s 21ms/step - loss: 1.331
2 - accuracy: 0.5624
Epoch 2/5
938/938 [=====] - 20s 21ms/step - loss: 0.275
9 - accuracy: 0.9167
Epoch 3/5
938/938 [=====] - 20s 21ms/step - loss: 0.133
6 - accuracy: 0.9598
Epoch 4/5
938/938 [=====] - 20s 21ms/step - loss: 0.088
3 - accuracy: 0.9736
Epoch 5/5
938/938 [=====] - 20s 21ms/step - loss: 0.068
1 - accuracy: 0.9791
```

```
Out[14]: <tensorflow.python.keras.callbacks.History at 0x1cf87a3e160>
```

```
In [15]: 1 test_loss, test_acc = model.evaluate(x_test, y_test)
          2 print(f"Test accuracy: {test_acc:.3f}")
          3
          4
```

```
313/313 [=====] - 1s 4ms/step - loss: 0.0557  
- accuracy: 0.9826  
Test accuracy: 0.983
```