

Stoneburner, Kurt

• DSC 650 - Week 06

• Assignment 6.2 - ConvNet Model that classifies images in the CIFAR digital dataset.

This is mostly copy/pasted code from a website I found online. I've added notes to model layers to help with my understanding and includes notes on keras activation functions. I built the final model based on an evaluation of the validation model. I imported the plotting functions from a previous exercise.

Reference: <https://keras.io/api/datasets/cifar10/> (<https://keras.io/api/datasets/cifar10/>)

CIFAR and Data augmentation https://stepup.ai/train_data_augmentation_keras/ (https://stepup.ai/train_data_augmentation_keras/)

Max-Pooling Explained: <https://analyticsindiamag.com/max-pooling-in-convolutional-neural-network-and-its-features/> (<https://analyticsindiamag.com/max-pooling-in-convolutional-neural-network-and-its-features/>)

Conv2D Official Documentation: https://keras.io/api/layers/convolution_layers/convolution2d/ (https://keras.io/api/layers/convolution_layers/convolution2d/)

```
In [1]: 1 """
2 #/** Enable Plaid-Ml GPU backend for Radeon Cards. (it works faster
3 from os import environ
4
5 environ["KERAS_BACKEND"] = "plaidml.keras.backend"
6
7 import tensorflow.keras
8 """
9 print()
```

```
In [2]: 1 import os
2 import sys
3 # /** Imports and Load Data
4 import matplotlib.pyplot as plt
5 import numpy as np
6 import pandas as pd
7 import scipy
8
9 #/** Use the whole window in the IPYNB editor
10 from IPython.core.display import display, HTML
11 display(HTML("<style>.container { width:100% !important; }</style>"))
12
13 #/** Maximize columns and rows displayed by pandas
14 pd.set_option('display.max_rows', 100)
15 pd.set_option('display.max_columns', None)
```

```

16
17 import tensorflow as tf
18
19 from tensorflow import keras
20 from tensorflow.keras import layers, models
21
22 from tensorflow.keras.preprocessing.image import ImageDataGenerator
23 from tensorflow.keras.utils import to_categorical
24

```

In [3]:

```

1  #####
2  Plot a Fitted Models History of Loss and Accuracy
3  #####
4  def plot_model_history(input_history, loss='loss', acc='accuracy'):
5      print(input_history.history.keys())
6      loss = input_history.history[list(input_history.history.keys())[0]]
7      acc = input_history.history[list(input_history.history.keys())[1]]
8
9
10
11     epochs = range(1, len(loss) + 1)
12     plt.plot(epochs, acc, "b", label="Training Accuracy")
13     plt.title("Training Accuracy\nAccuracy should go up")
14     plt.xlabel("Epochs")
15     plt.ylabel("Loss")
16     plt.legend()
17
18
19     epochs = range(1, len(loss) + 1)
20     plt.plot(epochs, loss, "bo", label="Training Loss")
21
22     plt.title("Training Loss \nLoss should go down")
23     plt.xlabel("Epochs")
24     plt.ylabel("Loss")
25     plt.legend()
26     plt.show()
27
28 #####
29 Plot a Fitted Models History Training and Validation Loss
30 #####
31 def plot_model_validation(input_history):
32
33     Assign Loss/Accuracy/Validation Loss/Validation Accuracy
34     Based on Dictionary Key Tuple position
35
36     loss = input_history.history[list(input_history.history.keys())[0]]
37     acc = input_history.history[list(input_history.history.keys())[1]]
38     val_loss = input_history.history[list(input_history.history.keys()
39     val_acc = input_history.history[list(input_history.history.keys()
40     #print(loss, acc, val_loss, val_acc)
41     epochs = range(1, len(loss) + 1)
42     plt.plot(epochs, loss, "bo", label="Training loss")
43     plt.plot(epochs, val_loss, "b", label="Validation loss")
44     plt.title("Training and validation loss")
45     plt.xlabel("Epochs")
46     plt.ylabel("Loss")

```

```

47     plt.legend()
48     plt.show()
49
50     #!/*** Plot the Validation Set Accuracy
51     plt.clf()
52     plt.plot(epochs, acc, "bo", label="Training accuracy")
53     plt.plot(epochs, val_acc, "b", label="Validation accuracy")
54     plt.title("Training and validation accuracy")
55     plt.xlabel("Epochs")
56     plt.ylabel("Accuracy")
57     plt.legend()
58     plt.show()
59
60     def visualize_data(images, categories, class_names):
61         fig = plt.figure(figsize=(14, 6))
62         fig.patch.set_facecolor('white')
63         for i in range(3 * 7):
64             plt.subplot(3, 7, i+1)
65             plt.xticks([])
66             plt.yticks([])
67             plt.imshow(images[i])
68             class_index = categories[i].argmax()
69             plt.xlabel(class_names[class_index])
70     plt.show()
71

```

Assignment 6.2a

Using section 5.2 in Deep Learning with Python as a guide, create a ConvNet model that classifies images CIFAR10 small images classification dataset. Do not use dropout or data-augmentation in this part.

Save the model, predictions, metrics, and validation plots in the dsc650/assignments /assignment06/results directory. If you are using JupyterHub, you can include those plots in your Jupyter notebook.

In [4]:

```

1
2
3     #!/*** Inputs reflects the shape of each individual piece of data.
4     #!/*** The MNIST is a 28x28 single channel image.
5     #!/*** The third is 1 channel. This is a greyscale image, therefore i
6     #!/*** See Link above for further explanation
7
8     #!/*** Conv2D: Filters defines the number of tensors at the layer. Ke
9     #!/*** Conv2D reduces the image size by (filter_size - 1) in
10    #!/*** MaxPooling2D: Is a form of feature reduction. In this case, Th
11    #!/*** pool-size value (2x2) is kept. With a pool value
12    #!/*** by 75%.
13    #!/*** At each stage of max pooling, the image gets sma
14    #!/*** for relationships between the reduced features.
15
16

```

Reference: <https://machinelearningknowledge.ai/keras-activation-layers-ultimate-guide-for-beginners/> (<https://machinelearningknowledge.ai/keras-activation-layers-ultimate-guide-for->

[beginners/\)](#)

Keras Activation Cheat Sheet

relu: Rectified linear unit.

- ReLu activation function is computationally efficient hence it enables neural networks to converge faster during the training phase.
- It is both non-linear and differentiable which are good characteristics for activation function.
- ReLU does not suffer from the issue of Vanishing Gradient issue like other activation functions and hence it is very effective in hidden layers of large neural networks.

Data used in **ReLU** layers must be positive. Negative numbers cannot be back-propagated and these values basically 'die'.

softmax: Generates a weighted value between 0 and 1. Use for multiple classification tasks.

sigmoid: Generates either 0 or 1. Data must be normalized. Does not handle outliers well Use for Binary Classification tasks.

tanh: Generates values from -1 to +1. Data must be normalized. Does not handle outliers well.

Which Activation Function to use in Neural Network?

Sigmoid and **Tanh** activation function should be avoided in hidden layers as they suffer from Vanishing Gradient problem.

Sigmoid activation function should be used in the output layer in case of Binary Classification

ReLU activation functions are ideal for hidden layers of neural networks as they do not suffer from the Vanishing Gradient problem and are computationally fast.

Softmax activation function should be used in the output layer in case of multiclass classification.

Label	Description
0	airplane
1	automobile
2	bird
3	cat
4	deer
5	dog
6	frog
7	horse
8	ship

Label Description

Returns

Tuple of NumPy arrays: (x_train, y_train), (x_test, y_test).

x_train: uint8 NumPy array of grayscale image data with shapes (50000, 32, 32, 3), containing the training data. Pixel values range from 0 to 255.

y_train: uint8 NumPy array of labels (integers in range 0-9) with shape (50000, 1) for the training data.

x_test: uint8 NumPy array of grayscale image data with shapes (10000, 32, 32, 3), containing the test data. Pixel values range from 0 to 255.

y_test: uint8 NumPy array of labels (integers in range 0-9) with shape (10000, 1) for the test data.

```
In [5]: 1 class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog']
2 num_classes = len(class_names)
3
4 #!/*** Load the CIFAR10 dataset into the default test train splits
5 (x_train, y_train), (x_test, y_test) = tf.keras.datasets.cifar10.load
6
7
8 #!/*** Build a subset of the data for model exploration
9 subset_size = 5000
10
11 #!/*** Generate Partial validation sets
12 x_validation = x_train[subset_size:subset_size*2]
13 y_validation = y_train[subset_size:subset_size*2]
14
15 #!/*** Validation subset
16 x_train = x_train[:subset_size]
17 y_train = y_train[:subset_size]
18
19
20 print(x_train.shape)
21 print(y_train.shape)
22
23 #!/*** Verify the subsets are the proper shape
24 assert x_train.shape == (subset_size, 32, 32, 3)
25 assert y_train.shape == (subset_size, 1)
26
27 #!/*** Max test size is 10000
28 if subset_size > 10000:
29     subset_size = 10000
30
31 x_test = x_test[:subset_size]
32 y_test = y_test[:subset_size]
33
34 print(x_test.shape)
35 print(y_test.shape)
36
37 #!/*** Verify the subsets are the proper shape
```

```

38 assert x_test.shape == (subset_size, 32, 32, 3)
39 assert y_test.shape == (subset_size, 1)
40
41
42 x_train = x_train / 255.0
43 y_train = to_categorical(y_train, num_classes)
44
45 x_test = x_test / 255.0
46 y_test = to_categorical(y_test, num_classes)
47

```

```

(5000, 32, 32, 3)
(5000, 1)
(5000, 32, 32, 3)
(5000, 1)

```



In [6]:

```

1 def create_model():
2
3     #!/** Using relu (Rectified Linear Units) for the hidden layers.
4     #!/** Using Convolution2d kernel_size =3, create unique features
5     #!/** MaxPooling 2, takes the 3x3 feature pixels (which are now
6     #!/** This reduces the features by 75% per feature group. This i
7
8     model = models.Sequential()
9     model.add(layers.Conv2D(32, (3, 3), activation='relu', padding='s
10    model.add(layers.Conv2D(32, (3, 3), activation='relu', padding='s
11    model.add(layers.MaxPool2D((2,2)))
12
13    model.add(layers.Conv2D(64, (3, 3), activation='relu', padding='s
14    model.add(layers.Conv2D(64, (3, 3), activation='relu', padding='s
15    model.add(layers.MaxPool2D((2,2)))
16
17    model.add(layers.Conv2D(128, (3, 3), activation='relu', padding='
18    model.add(layers.Conv2D(128, (3, 3), activation='relu', padding='
19    model.add(layers.MaxPool2D((2,2)))
20
21    model.add(layers.Flatten())
22    model.add(layers.Dense(128, activation='relu'))
23    model.add(layers.Dense(10, activation='softmax'))
24

```

```
25     model.compile(optimizer='adam',
26                   loss='categorical_crossentropy',
27                   metrics=['accuracy'])
28
29     return model
```

```
In [7]: 1 batch_size = 32
        2 epochs = 30
        3 m_no_aug = create_model()
        4 m_no_aug.summary()
        5
        6 history_no_aug = m_no_aug.fit(
        7     x_train, y_train,
        8     epochs=epochs, batch_size=batch_size,
        9     validation_data=(x_test, y_test))
        10
        11 loss_no_aug, acc_no_aug = m_no_aug.evaluate(x_test, y_test)
        12
        13 print(f"Test accuracy: {acc_no_aug:.3f}")
        14
        15
        16
        17 plot_model_validation(history_no_aug)
```

WARNING:tensorflow:From C:\Users\family\anaconda3\envs\directml\lib\site-packages\tensorflow_core\python\ops\resource_variable_ops.py:1630: calling BaseResourceVariable.__init__ (from tensorflow.python.ops.resource_variable_ops) with constraint is deprecated and will be removed in a future version.

Instructions for updating:

If using Keras pass *_constraint arguments to layers.

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	896
conv2d_1 (Conv2D)	(None, 32, 32, 32)	9248
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_2 (Conv2D)	(None, 16, 16, 64)	18496
conv2d_3 (Conv2D)	(None, 16, 16, 64)	36928
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 64)	0
conv2d_4 (Conv2D)	(None, 8, 8, 128)	73856
conv2d_5 (Conv2D)	(None, 8, 8, 128)	147584
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 128)	0
flatten (Flatten)	(None, 2048)	0
dense (Dense)	(None, 128)	262272
dense_1 (Dense)	(None, 10)	1290

Total params: 550,570

Trainable params: 550,570

Non-trainable params: 0

Train on 5000 samples, validate on 5000 samples

Epoch 1/30

5000/5000 [=====] - 4s 869us/sample - loss: 2.0975 - acc: 0.2126 - val_loss: 1.8855 - val_acc: 0.3104

Epoch 2/30

5000/5000 [=====] - 3s 657us/sample - loss: 1.7665 - acc: 0.3472 - val_loss: 1.7073 - val_acc: 0.3908

Epoch 3/30

5000/5000 [=====] - 3s 656us/sample - loss: 1.5784 - acc: 0.4224 - val_loss: 1.5296 - val_acc: 0.4474

Epoch 4/30

5000/5000 [=====] - 3s 660us/sample - loss: 1.4523 - acc: 0.4720 - val_loss: 1.5033 - val_acc: 0.4514


Epoch 5/30

5000/5000 [=====] - 3s 660us/sample - loss: 1.2895 - acc: 0.5338 - val_loss: 1.4685 - val_acc: 0.4896

Epoch 6/30


```
5000/5000 [=====] - 3s 661us/sample - loss:
1.1680 - acc: 0.5708 - val_loss: 1.3664 - val_acc: 0.5032
Epoch 7/30
5000/5000 [=====] - 3s 656us/sample - loss:
1.0337 - acc: 0.6328 - val_loss: 1.4326 - val_acc: 0.5174
Epoch 8/30
5000/5000 [=====] - 3s 654us/sample - loss:
0.8545 - acc: 0.6914 - val_loss: 1.5443 - val_acc: 0.4996
Epoch 9/30
5000/5000 [=====] - 3s 653us/sample - loss:
0.6645 - acc: 0.7600 - val_loss: 1.7166 - val_acc: 0.5122
Epoch 10/30
5000/5000 [=====] - 3s 656us/sample - loss:
0.5287 - acc: 0.8070 - val_loss: 1.7572 - val_acc: 0.5014
Epoch 11/30
5000/5000 [=====] - 3s 657us/sample - loss:
0.3674 - acc: 0.8742 - val_loss: 1.9354 - val_acc: 0.5068
Epoch 12/30
5000/5000 [=====] - 3s 653us/sample - loss:
0.2485 - acc: 0.9124 - val_loss: 2.2690 - val_acc: 0.5240
Epoch 13/30
5000/5000 [=====] - 3s 656us/sample - loss:
0.2210 - acc: 0.9268 - val_loss: 2.8940 - val_acc: 0.5268
Epoch 14/30
5000/5000 [=====] - 3s 660us/sample - loss:
0.1439 - acc: 0.9504 - val_loss: 3.0321 - val_acc: 0.5024
Epoch 15/30
5000/5000 [=====] - 3s 658us/sample - loss:
0.2157 - acc: 0.9272 - val_loss: 2.7528 - val_acc: 0.5126
Epoch 16/30
5000/5000 [=====] - 3s 661us/sample - loss:
0.1037 - acc: 0.9638 - val_loss: 3.0390 - val_acc: 0.5182
Epoch 17/30
5000/5000 [=====] - 3s 658us/sample - loss:
0.1160 - acc: 0.9610 - val_loss: 3.5346 - val_acc: 0.5128
Epoch 18/30
5000/5000 [=====] - 3s 654us/sample - loss:
0.1069 - acc: 0.9632 - val_loss: 3.4819 - val_acc: 0.5178
Epoch 19/30
5000/5000 [=====] - 3s 658us/sample - loss:
0.0992 - acc: 0.9670 - val_loss: 3.1979 - val_acc: 0.5090
Epoch 20/30
5000/5000 [=====] - 3s 656us/sample - loss:
```





Run No-Augmentation Model with full dataset

14 Epochs is seems to be a good balance between fitting and over fitting - Using 5000 Samples.

22 Epochs was a good value when testing with 1000 Samples.

In [8]:

```

1  #!/*** Reload full data set
2  class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog'
3  num_classes = len(class_names)
4
5  #!/*** Load the CIFAR10 dataset into the default test train splits
6  (x_train, y_train), (x_test, y_test) = tf.keras.datasets.cifar10.load
7
8
9  #!/*** Verify the subsets are the proper shape
10 assert x_train.shape == (50000, 32, 32, 3)
11 assert y_train.shape == (50000, 1)
12
13 #!/*** Verify the subsets are the proper shape
14 assert x_test.shape == (10000, 32, 32, 3)
15 assert y_test.shape == (10000, 1)
16
17 #!/*** Not sure why x_train is divided by 255 or why Y_train is conve
18 x_train = x_train / 255.0
19 y_train = to_categorical(y_train, num_classes)
20
21 x_test = x_test / 255.0
22 y_test = to_categorical(y_test, num_classes)
23
24 #!/*** Display Sample of Data
25 visualize_data(y_train, x_train, class_names)

```



In [9]:

```

1  #!/*** 14 Epochs to prevent overfitting
2  batch_size = 32
3  epochs = 14
4  m_no_aug = create_model()
5  m_no_aug.summary()
6
7  history_no_aug = m_no_aug.fit(
8      x_train, y_train,
9      epochs=epochs, batch_size=batch_size,
10     validation_data=(x_test, y_test))
11
12 loss_no_aug, acc_no_aug = m_no_aug.evaluate(x_test, y_test)
13
14 print(f"Model accuracy: {acc_no_aug:.3f}")

```

15

16

17 plot_model_history(history, no_avg)

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_6 (Conv2D)	(None, 32, 32, 32)	896
conv2d_7 (Conv2D)	(None, 32, 32, 32)	9248
max_pooling2d_3 (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_8 (Conv2D)	(None, 16, 16, 64)	18496
conv2d_9 (Conv2D)	(None, 16, 16, 64)	36928
max_pooling2d_4 (MaxPooling2D)	(None, 8, 8, 64)	0
conv2d_10 (Conv2D)	(None, 8, 8, 128)	73856
conv2d_11 (Conv2D)	(None, 8, 8, 128)	147584
max_pooling2d_5 (MaxPooling2D)	(None, 4, 4, 128)	0
flatten_1 (Flatten)	(None, 2048)	0
dense_2 (Dense)	(None, 128)	262272
dense_3 (Dense)	(None, 10)	1290
Total params: 550,570		
Trainable params: 550,570		
Non-trainable params: 0		

Train on 50000 samples, validate on 10000 samples

Epoch 1/14

50000/50000 [=====] - 26s 510us/sample - loss: 1.5360 - acc: 0.4358 - val_loss: 1.1374 - val_acc: 0.5962

Epoch 2/14

50000/50000 [=====] - 25s 506us/sample - loss: 1.0162 - acc: 0.6391 - val_loss: 0.9327 - val_acc: 0.6672

Epoch 3/14

50000/50000 [=====] - 25s 508us/sample - loss: 0.8172 - acc: 0.7136 - val_loss: 0.8484 - val_acc: 0.7025

Epoch 4/14

50000/50000 [=====] - 26s 512us/sample - loss: 0.6937 - acc: 0.7570 - val_loss: 0.8370 - val_acc: 0.7165

Epoch 5/14

50000/50000 [=====] - 26s 511us/sample - loss: 0.5950 - acc: 0.7911 - val_loss: 0.8198 - val_acc: 0.7286

Epoch 6/14

50000/50000 [=====] - 25s 507us/sample - loss: 0.5169 - acc: 0.8170 - val_loss: 0.7811 - val_acc: 0.7398


Epoch 7/14

50000/50000 [=====] - 25s 510us/sample - loss:

```

s: 0.4428 - acc: 0.8422 - val_loss: 0.8098 - val_acc: 0.7388
Epoch 8/14
50000/50000 [=====] - 25s 508us/sample - los
s: 0.3873 - acc: 0.8618 - val_loss: 0.8176 - val_acc: 0.7438
Epoch 9/14
50000/50000 [=====] - 25s 506us/sample - los
s: 0.3289 - acc: 0.8820 - val_loss: 0.8895 - val_acc: 0.7450
Epoch 10/14
50000/50000 [=====] - 25s 506us/sample - los
s: 0.2867 - acc: 0.8988 - val_loss: 0.9468 - val_acc: 0.7487
Epoch 11/14
50000/50000 [=====] - 25s 507us/sample - los
s: 0.2551 - acc: 0.9101 - val_loss: 1.0424 - val_acc: 0.7423
Epoch 12/14
50000/50000 [=====] - 25s 506us/sample - los
s: 0.2286 - acc: 0.9188 - val_loss: 1.1135 - val_acc: 0.7479
Epoch 13/14
50000/50000 [=====] - 25s 506us/sample - los
s: 0.2111 - acc: 0.9253 - val_loss: 1.1478 - val_acc: 0.7442
Epoch 14/14
50000/50000 [=====] - 25s 506us/sample - los
s: 0.1947 - acc: 0.9312 - val_loss: 1.1876 - val_acc: 0.7382
10000/10000 [=====] - 2s 191us/sample - loss:
1.1876 - acc: 0.7382
Model accuracy: 0.738
dict keys(['loss', 'acc', 'val loss', 'val acc'])

```



Assignment 6.2.b

Using section 5.2 in Deep Learning with Python as a guide, create a ConvNet model that classifies images CIFAR10 small images classification dataset. This time includes dropout and data-augmentation.

```

In [10]: 1 import scipy
          2 class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog'
          3 num_classes = len(class_names)

```


```

4
5 #!/*** Load the CIFAR10 dataset into the default test train splits
6 (x_train, y_train), (x_test, y_test) = tf.keras.datasets.cifar10.load
7
8 x_train = x_train / 255.0
9 y_train = to_categorical(y_train, num_classes)
10
11 x_test = x_test / 255.0
12 y_test = to_categorical(y_test, num_classes)
13
14 visualize_data(x_train, y_train, class_names)
15
16 width_shift = 3/32
17 height_shift = 3/32
18 flip = True
19
20 datagen = ImageDataGenerator(
21     horizontal_flip=flip,
22     width_shift_range=width_shift,
23     height_shift_range=height_shift,
24 )
25 datagen.fit(x_train)
26
27 it = datagen.flow(x_train, y_train, shuffle=False)
28 batch_images, batch_labels = next(it)
29 print("Transformed images")
30 visualize_data(batch_images, batch_labels, class_names)

```



Transformed images



In [11]:

```
1 batch_size = 32
2 epochs = 14
3 model = create_model()
4 model.summary()
5
6 history = model.fit(
7     x_train, y_train,
8     epochs=epochs, batch_size=batch_size,
9     #validation_data=(x_test, y_test)
10 )
11
12 loss, acc = model.evaluate(batch_images, batch_labels)
13
14 print(f"Model accuracy: {acc:.3f}")
15
16
17 plot_model_history(history)
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
conv2d_12 (Conv2D)	(None, 32, 32, 32)	896
conv2d_13 (Conv2D)	(None, 32, 32, 32)	9248
max_pooling2d_6 (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_14 (Conv2D)	(None, 16, 16, 64)	18496
conv2d_15 (Conv2D)	(None, 16, 16, 64)	36928
max_pooling2d_7 (MaxPooling2D)	(None, 8, 8, 64)	0
conv2d_16 (Conv2D)	(None, 8, 8, 128)	73856
conv2d_17 (Conv2D)	(None, 8, 8, 128)	147584
max_pooling2d_8 (MaxPooling2D)	(None, 4, 4, 128)	0
flatten_2 (Flatten)	(None, 2048)	0
dense_4 (Dense)	(None, 128)	262272
dense_5 (Dense)	(None, 10)	1290
Total params: 550,570		
Trainable params: 550,570		
Non-trainable params: 0		

Train on 50000 samples

Epoch 1/14

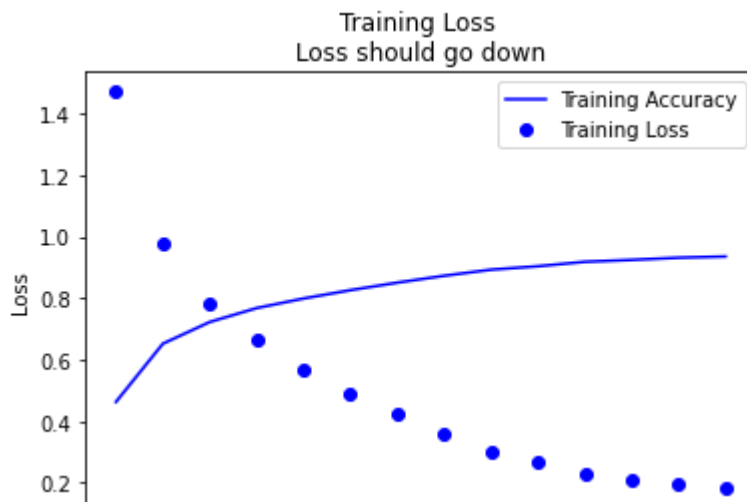
50000/50000 [=====] - 24s 470us/sample - loss: 1.4707 - acc: 0.4629

Epoch 2/14

50000/50000 [=====] - 23s 467us/sample - loss: 0.9800 - acc: 0.6520

Epoch 3/14

50000/50000 [=====] - 23s 468us/sample - loss:



In []:

In []:

In []: