

Assignment 9.2

By Kurt Stoneburner

Pyspark Time/Session Windows: <https://towardsdatascience.com/spark-3-2-session-windowing-feature-for-streaming-data-e404d92e267> (<https://towardsdatascience.com/spark-3-2-session-windowing-feature-for-streaming-data-e404d92e267>)

```
In [1]: 1 import os
2 import shutil
3 import json
4 from pathlib import Path
5 from IPython.display import clear_output
6 import pandas as pd
7
8 from kafka import KafkaProducer, KafkaAdminClient
9 from kafka.admin.new_topic import NewTopic
10 from kafka.errors import TopicAlreadyExistsError
11
12 from pyspark.sql import SparkSession
13 from pyspark.streaming import StreamingContext
14 from pyspark import SparkConf
15 from pyspark.sql.functions import window, from_json, col
16 from pyspark.sql.types import StringType, TimestampType, DoubleType,
17 from pyspark.sql.functions import udf
18
19 current_dir = Path(os.getcwd()).absolute()
20 checkpoint_dir = current_dir.joinpath('checkpoints')
21 locations_windowed_checkpoint_dir = checkpoint_dir.joinpath('location
22
23 if locations_windowed_checkpoint_dir.exists():
24     shutil.rmtree(locations_windowed_checkpoint_dir)
25
26
```

Configuration Parameters

TODO: Change the configuration parameters to the appropriate values for your setup.

```
In [2]: 1 config = dict(
2     bootstrap_servers=['kafka.kafka.svc.cluster.local:9092'],
3     first_name='Kurt',
4     last_name='Stoneburner'
5 )
6
7 config['client_id'] = '{}{}'.format(
8     config['last_name'],
9     config['first_name']
```

```

10 )
11 config['topic_prefix'] = '{}{}'.format(
12     config['last_name'],
13     config['first_name']
14 )
15
16 config['locations_topic'] = '{}-locations'.format(config['topic_prefix'])
17 config['accelerations_topic'] = '{}-accelerations'.format(config['topic_prefix'])
18 config['windowed_topic'] = '{}-windowed'.format(config['topic_prefix'])
19
20

```

```

Out[2]: {'bootstrap_servers': ['kafka.kafka.svc.cluster.local:9092'],
        'first_name': 'Kurt',
        'last_name': 'Stoneburner',
        'client_id': 'StoneburnerKurt',
        'topic_prefix': 'StoneburnerKurt',
        'locations_topic': 'StoneburnerKurt-locations',
        'accelerations_topic': 'StoneburnerKurt-accelerations',
        'windowed_topic': 'StoneburnerKurt-windowed'}

```

Create Topic Utility Function

The `create_kafka_topic` helps create a Kafka topic based on your configuration settings.

For instance, if your first name is *John* and your last name is *Doe*,

`create_kafka_topic('locations')` will create a topic with the name `DoeJohn-locations`. The function will not create the topic if it already exists.

```

In [3]: 1 def create_kafka_topic(topic_name, config=config, num_partitions=1, replication_factor=1):
2         bootstrap_servers = config['bootstrap_servers']
3         client_id = config['client_id']
4         topic_prefix = config['topic_prefix']
5         name = '{}-{}'.format(topic_prefix, topic_name)
6
7         admin_client = KafkaAdminClient(
8             bootstrap_servers=bootstrap_servers,
9             client_id=client_id
10        )
11
12        topic = NewTopic(
13            name=name,
14            num_partitions=num_partitions,
15            replication_factor=replication_factor
16        )
17
18        topic_list = [topic]
19        try:
20            admin_client.create_topics(new_topics=topic_list)
21            print('Created topic "{}"'.format(name))
22        except TopicAlreadyExistsError as e:
23            print('Topic "{}" already exists'.format(name))
24
25        create_kafka_topic('windowed')
Topic "StoneburnerKurt-windowed" already exists

```

TODO: This code is identical to the code used in 9.1 to publish acceleration and location data to the `LastnameFirstname-simple` topic. You will need to add in the code you used to create the `df_accelerations` dataframe. In order to read data from this topic, make sure that you are running the notebook you created in assignment 8 that publishes acceleration and location data to the `LastnameFirstname-simple` topic.

```
In [4]: 1 spark = SparkSession\
2         .builder\
3         .appName("Assignment09")\
4         .getOrCreate()
5
6 df_locations = spark \
7     .readStream \
8     .format("kafka") \
9     .option("kafka.bootstrap.servers", config['bootstrap_servers'][0])
10    .option("subscribe", config['locations_topic']) \
11    .load()
12
13 df_accelerations = spark\
14     .readStream.format("kafka")\
15     .option("kafka.bootstrap.servers", config['bootstrap_servers'][0])
16     .option("subscribe", config['accelerations_topic'])\
17     .load()
```

The following code defines a Spark schema for location and acceleration data as well as a user-defined function (UDF) for parsing the location and acceleration JSON data.

```
In [5]: 1 location_schema = StructType([
2         StructField('offset', DoubleType(), nullable=True),
3         StructField('id', StringType(), nullable=True),
4         StructField('ride_id', StringType(), nullable=True),
5         StructField('uuid', StringType(), nullable=True),
6         StructField('course', DoubleType(), nullable=True),
7         StructField('latitude', DoubleType(), nullable=True),
8         StructField('longitude', DoubleType(), nullable=True),
9         StructField('geohash', StringType(), nullable=True),
10        StructField('speed', DoubleType(), nullable=True),
11        StructField('accuracy', DoubleType(), nullable=True),
12    ])
13
14 acceleration_schema = StructType([
15     StructField('offset', DoubleType(), nullable=True),
16     StructField('id', StringType(), nullable=True),
17     StructField('ride_id', StringType(), nullable=True),
18     StructField('uuid', StringType(), nullable=True),
19     StructField('x', DoubleType(), nullable=True),
20     StructField('y', DoubleType(), nullable=True),
21     StructField('z', DoubleType(), nullable=True),
22 ])
23
24 udf_parse_acceleration = udf(lambda x: json.loads(x.decode('utf-8')),
25 udf_parse_location = udf(lambda x: json.loads(x.decode('utf-8')), loc
```

See <http://spark.apache.org/docs/latest/structured-streaming-programming-guide.html#window->

[operations-on-event-time \(http://spark.apache.org/docs/latest/structured-streaming-programming-guide.html#window-operations-on-event-time\)](http://spark.apache.org/docs/latest/structured-streaming-programming-guide.html#window-operations-on-event-time) for details on how to implement windowed operations.

The following code selects the `timestamp` column from the `df_locations` dataframe that reads from the `LastnameFirstname-locations` topic and parses the binary value using the `udf_parse_location` UDF and defines the result to the `json_value` column.

```
df_locations \
    .select(
        col('timestamp'),
        udf_parse_location(df_locations['value']).alias('json_value')
    )
```

From here, you can select data from the `json_value` column using the `select` method. For instance, if you saved the results of the previous code snippet to `df_locations_parsed` you could select columns from the `json_value` field and assign them aliases using the following code.

```
df_locations_parsed.select(
    col('timestamp'),
    col('json_value.ride_id').alias('ride_id'),
    col('json_value.uuid').alias('uuid'),
    col('json_value.speed').alias('speed')
)
```

Next, you will want to add a watermark and group by `ride_id` and `speed` using a window duration of *30 seconds* and a slide duration of *15 seconds*. Use the `withWatermark` method in conjunction with the `groupBy` method. The [Spark streaming documentation \(http://spark.apache.org/docs/latest/structured-streaming-programming-guide.html#window-operations-on-event-time\)](http://spark.apache.org/docs/latest/structured-streaming-programming-guide.html#window-operations-on-event-time) should provide examples of how to do this.

Next use the `mean` aggregation method to compute the average values and rename the column `avg(speed)` to `value` and the column `ride_id` to `key`. The reason you are renaming these values is that the PySpark Kafka API expects `key` and `value` as inputs. In a production example, you would setup serialization that would handle these details for you.

When you are finished, you should have a streaming query with `key` and `value` as columns.

```
In [6]: 1 df_locations_parsed = df_locations \
2         .select(
3             col('timestamp'),
4             udf_parse_location(df_locations['value']).alias('json_value')
5         )
6
7 df_locations_parsed.select(
8     col('timestamp'),
9     col('json_value.ride_id').alias('ride_id'),
10    col('json_value.uuid').alias('uuid'),
11    col('json_value.speed').alias('speed')
12 )
```

```

13
14
15
root
|-- timestamp: timestamp (nullable = true)
|-- json_value: struct (nullable = true)
|   |-- offset: double (nullable = true)
|   |-- id: string (nullable = true)
|   |-- ride_id: string (nullable = true)
|   |-- uuid: string (nullable = true)
|   |-- course: double (nullable = true)
|   |-- latitude: double (nullable = true)
|   |-- longitude: double (nullable = true)
|   |-- geohash: string (nullable = true)
|   |-- speed: double (nullable = true)
|   |-- accuracy: double (nullable = true)

```

```

In [7]: 1 """
2 # Group the data by session window and userId, and compute the count
3
4         sessionizedCounts = events \
5             .withWatermark("timestamp", "10 minutes") \
6             .groupBy(
7                 session_window(events.timestamp, "5 minutes"),
8                 events.userId) \
9             .count()
10
11 slidingWindows = windowing_df.withWatermark("timeReceived", "10 minutes") \
12     .groupBy("eventId", window("timeReceived", "10 minutes", "5 minutes")) \
13     .count()
14 """

```

```

Out[7]: '\n# Group the data by session window and userId, and compute the count of each group\n\n    sessionizedCounts = events        .withWatermark("timestamp", "10 minutes")        .groupBy(\n        session_window(event\ns.timestamp, "5 minutes"),\n        events.userId)        .count()\n\nslidingWindows = windowing_df.withWatermark("timeReceived", "10 minutes")\n\nslidingWindows.groupBy("eventId", window("timeReceived", "10 minutes", "5 minutes"))\n\nslidingWindows.count()\n\nslidingWindows.show(truncate = False)\n\n'

```

```

In [8]: 1 windowedSpeeds = df_locations_parsed \
2         .withWatermark("timestamp", "10 seconds") \
3         .groupBy("json_value.ride_id", window("timestamp", "10 seconds",
4         .mean('json_value.speed') \
5         .withColumnRenamed("ride_id", "key") \
6         .withColumnRenamed("avg(json_value.speed AS `speed`)", "value") \
7         .select(col('key'), col('value'))
8
9
root
|-- key: string (nullable = true)
|-- value: double (nullable = true)

```

```

In [9]: 1 windowedSpeeds = df_locations_parsed \
2         .withWatermark("timestamp", "10 seconds") \
3         .groupBy("json_value.ride_id", window("timestamp", "10 seconds",

```

```

4     .mean('json_value.speed') \
5     .withColumnRenamed("ride_id", "key") \
6     .withColumnRenamed("avg(json_value.speed AS `speed`)", "value") \
7     .select(col('key'), col('value'))
8
root
|-- key: string (nullable = true)
|-- value: double (nullable = true)

```

```

In [10]: 1 windowedSpeeds = df_locations_parsed \
2         .withWatermark("timestamp", "10 seconds") \
3         .groupBy(window("timestamp", "10 seconds", "5 seconds"), "json_val
4         .count() \
5         .select(col('ride_id').alias("key"), col('count').alias("value"))
6
7
root
|-- key: string (nullable = true)
|-- value: long (nullable = false)

```

```

In [11]: 1 """
2 windowedSpeeds = df_locations_parsed \
3         .withWatermark('timestamp', "30 seconds") \
4         .groupBy(
5             window("timestamp", "30 seconds", "15 seconds"),
6             "json_value.ride_id") \
7         .mean("json_value.speed")
8
9 windowedSpeeds.printSchema()
10 windowedSpeeds.select(windowedSpeeds.columns[2])
11
12 #.select(col("json_value.ride_id").alias("key"), col("avg(speed)").al
13
14 """

```

```

Out[11]: '\nwindowedSpeeds = df_locations_parsed        .withWatermark(\'timestam
p\', "30 seconds")        .groupBy(\n                window("timestamp", "30 se
conds", "15 seconds"),\n                "json_value.ride_id")        .mean("jso
n_value.speed") \n        \nwindowedSpeeds.printSchema()\nwindowedSpeeds.s
elect(windowedSpeeds.columns[2])\n\n#.#.select(col("json_value.ride_i
d").alias("key"), col("avg(speed)").alias("value"))\n\n'

```

```

In [12]: 1 """
2 def foreach_batch_function(df, epoch_id):
3
4     # Transform and write batchDF
5     #print(df.printSchema())
6     try:
7         clear_output(wait=True)
8         print(df.select(df.columns).show())
9     except:
10         pass

```

```

11
12
13
14 ds_locations = df_locations_parsed.writeStream.foreachBatch(foreach_b
15
16
17
18 try:
19     ds_locations.awaitTermination()
20 except KeyboardInterrupt:
21     print("STOPPING STREAMING DATA")
22 """
23

```

```

In [13]: 1 windowedSpeeds = df_locations_parsed \
2         .withWatermark("timestamp", "15 seconds") \
3         .groupBy(
4             window("timestamp", "30 seconds", "15 seconds"),
5             "json_value.ride_id",
6             'json_value.speed',
7             'timestamp'
8         ) \
9         .mean('json_value.speed') \
10        .select(col('ride_id').alias('key'), col('speed').alias('value'))
11
12
13
14
15
16
17
18 root
19 |-- key: string (nullable = true)
20 |-- value: double (nullable = true)
21

```

In []:

```
1
```

In the previous Jupyter cells, you should have created the `windowedSpeeds` streaming query. Next, you will need to write that to the `LastnameFirstname-windowed` topic. If you created the `windowsSpeeds` streaming query correctly, the following should publish the results to the `LastnameFirstname-windowed` topic.

```

In [21]: 1 ds_locations_windowed = windowedSpeeds \
2         .selectExpr("CAST(key AS STRING)", "CAST(value AS STRING)") \
3         .writeStream \
4         .format("kafka") \
5         .option("kafka.bootstrap.servers", "kafka.kafka.svc.cluster.local:9
6         .option("topic", config['windowed_topic']) \
7         .option("checkpointLocation", str(locations_windowed_checkpoint_dir
8         .start()
9
10 try:
11     ds_locations_windowed.awaitTermination()
12 except KeyboardInterrupt:
13     print("STOPPING STREAMING DATA")

```

```

-----
-----
StreamingQueryException                                Traceback (most recent call
last)
<ipython-input-21-0e36252d7665> in <module>
      9
     10 try:
--> 11     ds_locations_windowed.awaitTermination()
     12 except KeyboardInterrupt:
     13     print("STOPPING STREAMING DATA")

/usr/local/spark/python/pyspark/sql/streaming.py in awaitTermination(self, timeout)
     101         return self._jsq.awaitTermination(int(timeout * 10
00))
     102     else:
--> 103         return self._jsq.awaitTermination()
     104
     105     @property

/usr/local/spark/python/lib/py4j-0.10.9-src.zip/py4j/java_gateway.py in
n __call__(self, *args)
     1302
     1303         answer = self.gateway_client.send_command(command)
-> 1304         return_value = get_return_value(
     1305             answer, self.gateway_client, self.target_id, self.
name)
     1306

/usr/local/spark/python/pyspark/sql/utils.py in deco(*a, **kw)
     135         # Hide where the exception came from that show
s a non-Pythonic
     136         # JVM exception message.
--> 137         raise_from(converted)
     138     else:
     139         raise

/usr/local/spark/python/pyspark/sql/utils.py in raise_from(e)

StreamingQueryException: Writing job aborted.
=== Streaming Query ===
Identifier: [id = 60f34d1c-fc52-4f84-aeec-cdbf62465592, runId = 65cb0f
a9-3da1-45a4-ba24-a59c88991847]
Current Committed Offsets: {KafkaV2[Subscribe[StoneburnerKurt-location
s]]: {"StoneburnerKurt-locations":{"0":11104}}}
Current Available Offsets: {KafkaV2[Subscribe[StoneburnerKurt-location
s]]: {"StoneburnerKurt-locations":{"0":11105}}}

Current State: ACTIVE
Thread State: RUNNABLE

Logical Plan:
WriteToMicroBatchDataSource org.apache.spark.sql.kafka010.KafkaStreami
ngWrite@6b89d9cc
+- Project [cast(key#144 as string) AS key#636, cast(value#145 as stri
ng) AS value#637]

```



```

+- Project [ride_id#136 AS key#144, speed#137 AS value#145]
+- Aggregate [window#138-T15000ms, json_value#43.ride_id, json_v
alue#43.speed, timestamp#12-T15000ms], [window#138-T15000ms AS window#
128-T15000ms, json_value#43.ride_id AS ride_id#136, json_value#43.spee
d AS speed#137, timestamp#12-T15000ms, avg(json_value#43.speed) AS avg
(json_value.speed AS `speed`)#133]
+- Filter ((timestamp#12-T15000ms >= window#138-T15000ms.star
t) AND (timestamp#12-T15000ms < window#138-T15000ms.end))
+- Expand [ArrayBuffer(named_struct(start, precisetimestamp
pconversion((((CASE WHEN (cast(CEIL((cast((precisetimestampconversion
(timestamp#12-T15000ms, TimestampType, LongType) - 0) as double) / cas
t(15000000 as double))) as double) = (cast((precisetimestampconversion
(timestamp#12-T15000ms, TimestampType, LongType) - 0) as double) / cas
t(15000000 as double))) THEN (CEIL((cast((precisetimestampconversion(t
imestamp#12-T15000ms, TimestampType, LongType) - 0) as double) / cast
(15000000 as double))) + cast(1 as bigint)) ELSE CEIL((cast((preciseti
mestampconversion(timestamp#12-T15000ms, TimestampType, LongType) - 0)
as double) / cast(15000000 as double))) END + cast(0 as bigint)) - cas
t(2 as bigint)) * 15000000) + 0), LongType, TimestampType), end, preci
setimestampconversion((((CASE WHEN (cast(CEIL((cast((precisetimestamp
pconversion(timestamp#12-T15000ms, TimestampType, LongType) - 0) as do
uble) / cast(15000000 as double))) as double) = (cast((precisetimestam
pconversion(timestamp#12-T15000ms, TimestampType, LongType) - 0) as do
uble) / cast(15000000 as double))) THEN (CEIL((cast((precisetimestam
pconversion(timestamp#12-T15000ms, TimestampType, LongType) - 0) as doub
le) / cast(15000000 as double))) + cast(1 as bigint)) ELSE CEIL((cast
((precisetimestampconversion(timestamp#12-T15000ms, TimestampType, Lon
gType) - 0) as double) / cast(15000000 as double))) END + cast(0 as bi
gint)) - cast(2 as bigint)) * 15000000) + 0) + 30000000), LongType, Ti
mestampType)), timestamp#12-T15000ms, json_value#43), ArrayBuffer(name
d_struct(start, precisetimestampconversion((((CASE WHEN (cast(CEIL((c
ast((precisetimestampconversion(timestamp#12-T15000ms, TimestampType,
LongType) - 0) as double) / cast(15000000 as double))) as double) = (c
ast((precisetimestampconversion(timestamp#12-T15000ms, TimestampType,
LongType) - 0) as double) / cast(15000000 as double))) THEN (CEIL((cas
t((precisetimestampconversion(timestamp#12-T15000ms, TimestampType, Lo
ngType) - 0) as double) / cast(15000000 as double))) + cast(1 as begin
t)) ELSE CEIL((cast((precisetimestampconversion(timestamp#12-T15000ms,
TimestampType, LongType) - 0) as double) / cast(15000000 as double)))
END + cast(1 as bigint)) - cast(2 as bigint)) * 15000000) + 0), LongTy
pe, TimestampType), end, precisetimestampconversion((((CASE WHEN (ca
st(CEIL((cast((precisetimestampconversion(timestamp#12-T15000ms, Times
tampType, LongType) - 0) as double) / cast(15000000 as double))) as do
uble) = (cast((precisetimestampconversion(timestamp#12-T15000ms, Times
tampType, LongType) - 0) as double) / cast(15000000 as double))) THEN
(CEIL((cast((precisetimestampconversion(timestamp#12-T15000ms, Timesta
mpType, LongType) - 0) as double) / cast(15000000 as double))) + cast
(1 as bigint)) ELSE CEIL((cast((precisetimestampconversion(timestamp#1
2-T15000ms, TimestampType, LongType) - 0) as double) / cast(15000000 a
s double))) END + cast(1 as bigint)) - cast(2 as bigint)) * 15000000)
+ 0) + 30000000), LongType, TimestampType)), timestamp#12-T15000ms, js
on_value#43)], [window#138-T15000ms, timestamp#12-T15000ms, json_value
#43]

+- EventTimeWatermark timestamp#12: timestamp, 15 secon
ds

+- Project [timestamp#12, <lambda>(value#8) AS json_

```

```
value#43]
      +- StreamingDataSourceV2Relation [key#7, value#8,
topic#9, partition#10, offset#11L, timestamp#12, timestampType#13], or
g.apache.spark.sql.kafka010.KafkaSourceProvider$KafkaScan@31a12753, Ka
```