

Stoneburner, Kurt

. DSC 650 - Assignment 10

Links to Deep Learning Sample Code: Word Embedding Examples:

- https://github.com/fchollet/deep-learning-with-python-notebooks/blob/master/first_edition/6.1-using-word-embeddings.ipynb (https://github.com/fchollet/deep-learning-with-python-notebooks/blob/master/first_edition/6.1-using-word-embeddings.ipynb)

RNN and LSTM Examples

- https://github.com/fchollet/deep-learning-with-python-notebooks/blob/master/first_edition/6.2-understanding-recurrent-neural-networks.ipynb (https://github.com/fchollet/deep-learning-with-python-notebooks/blob/master/first_edition/6.2-understanding-recurrent-neural-networks.ipynb)

ngram reference:

- <https://www.analyticsvidhya.com/blog/2021/09/what-are-n-grams-and-how-to-implement-them-in-python/> (<https://www.analyticsvidhya.com/blog/2021/09/what-are-n-grams-and-how-to-implement-them-in-python/>)

pad_sequence: Used to transform lists within lists (2D array) to have a uniform inner dimension. Essentially, padding smaller arrays to the size of the largest, or trimming all arrays to max len

This function transforms a list (of length num_samples) of sequences (lists of integers) into a 2D Numpy array of shape (num_samples, num_timesteps). num_timesteps is either the maxlen argument if provided, or the length of the longest sequence in the list.

- https://www.tensorflow.org/api_docs/python/tf/keras/preprocessing/sequence/pad_sequences (https://www.tensorflow.org/api_docs/python/tf/keras/preprocessing/sequence/pad_sequences)

Working with tensors was a lost cause for this assignment. Keeping these for reference Convert Numpy Array to Tensor:

- <https://www.projectpro.io/recipes/convert-numpy-array-tensor> (<https://www.projectpro.io/recipes/convert-numpy-array-tensor>)

Convert Tensor to Numpy Array:

- <https://www.delftstack.com/howto/numpy/python-convert-tensor-to-numpy-array/> (<https://www.delftstack.com/howto/numpy/python-convert-tensor-to-numpy-array/>)

Basic Text Preparation for modeling: Tokenize, ngram, Convert Text to numeric categorical (integer) value. Use Keras pad_sequences to convert integer lists to uniform lists. Feed into model.

```

In [1]: 1 import os
        2 from pathlib import Path
        3 import sys
        4 # Imports and Load Data
        5 import matplotlib.pyplot as plt
        6 import numpy as np
        7 import pandas as pd
        8 import time
        9
       10 from tensorflow import keras
       11 import tensorflow as tf
       12 import datetime
       13 from tensorflow.keras.optimizers import RMSprop
       14 from tensorflow.keras import preprocessing
       15 from tensorflow.keras.models import Sequential
       16 from tensorflow.keras.layers import Flatten, Dense
       17 from tensorflow.keras.layers import Embedding
       18
       19
       20 Reusing Code from assignment 04
       21 from chardet.universaldetector import UniversalDetector
       22 from bs4 import BeautifulSoup
       23
       24
       25 import re
       26
       27 Use the whole window in the IPYNB editor
       28 from IPython.core.display import display, HTML
       29 display(HTML("<style>.container { width:100% !important; }</style>"))
       30
       31 Maximize columns and rows displayed by pandas
       32 pd.set_option('display.max_rows', 100)
       33 pd.set_option('display.max_columns', None)

```

```

In [2]: 1 *****
        2 Plot a Fitted Models History of Loss and Accuracy
        3 *****
        4 def plot_model_history(input_history):
        5     loss_key, acc_key = list(input_history.history.keys())[:2]
        6     val_loss_key, val_acc_key = list(input_history.history.keys())[2:
        7
        8     acc = input_history.history[acc_key]
        9     loss = input_history.history[loss_key]
       10
       11     val_loss = input_history.history[val_loss_key]
       12     val_acc = input_history.history[val_acc_key]
       13
       14     epochs = range(1, len(loss) + 1)
       15     plt.plot(epochs, acc, "b", label="Training Accuracy")
       16     plt.title("Training Accuracy\nAccuracy should go up")
       17     plt.xlabel("Epochs")
       18     plt.ylabel("Loss")
       19     plt.legend()
       20     plt.show()
       21

```

```
22     plt.plot(epochs, loss, "bo", label="Training Loss")
23
24     plt.title("Training Loss \nLoss should go down")
25     plt.xlabel("Epochs")
26     plt.ylabel("Loss")
27     plt.legend()
28     plt.show()
29
30
31     plt.plot(epochs, loss, "bo", label="Training loss")
32     plt.plot(epochs, val_loss, "b", label="Validation loss")
33     plt.title("Training and validation loss")
34     plt.xlabel("Epochs")
35     plt.ylabel("Loss")
36     plt.legend()
37     plt.show()
38
39     #!/*** Plot the Validation Set Accuracy
40     plt.clf()
41
42     plt.plot(epochs, acc, "bo", label="Training accuracy")
43     plt.plot(epochs, val_acc, "b", label="Validation accuracy")
44     plt.title("Training and validation accuracy")
45     plt.xlabel("Epochs")
46     plt.ylabel("Accuracy")
47     plt.legend()
48     plt.show()
49
50
```

In [3]:

```
1  #!/*** Get Working Directory
2  current_dir = Path(os.getcwd()).absolute()
3
4  #!/*** Go up Two folders
5  project_dir = current_dir.parents[2]
6
7  #!/*** IMDB Data Path
8  imdb_path = project_dir.joinpath("dsc650/data/external/imdb/aclImdb")
9
10 file_path = imdb_path.joinpath("train/pos")
11
12 #!/*** Grab the first positive review text for testing
13 file_path = file_path.joinpath(os.listdir(file_path)[0])
14
15 with open(file_path, 'r') as f:
16     sample_text = f.read()
17
18 print(sample_text)
19
```

Bromwell High is a cartoon comedy. It ran at the same time as some other programs about school life, such as "Teachers". My 35 years in the teaching profession lead me to believe that Bromwell High's satire is much closer to reality than is "Teachers". The scramble to survive fin

```
In [4]: 1  #!/*** Randomly assign 20% of the training Data and copy to a validat
2  import os, pathlib, shutil, random
3
4  val_dir = imdb_path.joinpath("val")
5  train_dir = imdb_path.joinpath("train")
6  test_dir = imdb_path.joinpath("test")
7
8  for category in ("neg", "pos"):
9      #!/*** Skip if val folder exists (Delete Folder to resample)
10     if os.path.exists(val_dir.joinpath(category)):
11         break
12
13     os.makedirs(val_dir.joinpath(category))
14     files = os.listdir(train_dir.joinpath(category))
15     random.Random(1337).shuffle(files)
16     num_val_samples = int(0.2 * len(files))
17     val_files = files[-num_val_samples:]
18     for fname in val_files:
19         shutil.move(train_dir / category / fname,
20                     val_dir / category / fname)
21
22
```

Load IMDB Dataset

```
In [5]: 1  #!/*** Use Universal Detector to determine file encoding.
2  #!/*** Borrowed from Assignment04
3  def read_file_with_encoding(filepath):
4
5      detector = UniversalDetector()
6
7      try:
8          with open(filepath) as f:
9              return f.read()
10     except UnicodeDecodeError:
11         detector.reset()
12         with open(filepath, 'rb') as f:
13             for line in f.readlines():
14                 detector.feed(line)
15                 if detector.done:
16                     break
17         detector.close()
18         encoding = detector.result['encoding']
19         with open(filepath, encoding=encoding) as f:
20             return f.read()
21
22     #!/*** Borrowed from Assignment04
23     def parse_html_payload(payload):
24         """
25         This function uses BeautifulSoup to read HTML data
```

```

26         and return the text.  If the payload is plain text, then
27         BeautifulSoup will return the original content
28         """
29         soup = BeautifulSoup(payload, 'html.parser')
30         return str(soup.get_text()).encode('utf-8').decode('utf-8')
31
32     def load_dataset(dir_path):
33
34         text = []
35         targets = []
36
37         #!/*** Crawl the neg and pos folders
38         for category in ("neg", "pos"):
39             files = os.listdir(dir_path.joinpath(category))
40
41             #!/*** Loop through each file in the folder
42             for file in files:
43                 try:
44                     #!/*** Add processed file to text
45                     text.append(
46                         #!/*** Strip HTML Tags
47                         parse_html_payload(
48                             #!/*** Read File from disk. Function uses Uni
49                             read_file_with_encoding(
50                                 dir_path.joinpath(category).joinpath(file
51
52                             #!/*** Append Target Value
53                             if category == 'neg':
54                                 targets.append(0)
55                             else:
56                                 targets.append(1)
57                 except:
58                     print(f"Dropping File: {file} due to decoding issues")
59
60         #!/*** Targets returned as Numpy float32 array
61         return text, np.asarray(targets).astype('float32')

```

Assignment 10.1

In [27]:

```

1  #!/*** Vectorize a corpus
2  class Vectorizer:
3      def __init__(self, **kwargs):
4          self.corpus_tokens = []
5          self.corpus_ngrams = []
6
7          self.max_tokens = None
8          self.ngram_size = 1
9          self.tidyup = True
10
11         self.max_element_count = -1
12
13         for key, value in kwargs.items():
14             if key == "max_tokens":
15                 self.max_tokens = value

```

```
16
17         if key == "ngrams":
18             self.ngram_size = value
19
20         if key == "tidyup":
21             self.tidyup = value
22
23
24     #!/*** One Hot Encoding Dictionaries
25     #!/*** Key = Token Index, Value = Word
26     self.ngram_index = {}
27
28     #!/*** Key = Word, Value = Token Index
29     self.vocabulary_index = {}
30
31     def tokenize(self, raw_text):
32         #!/*** Initialize Output Tokens
33         tokens = []
34
35         #!/*** Split Text into words
36         for x in re.split("\s", raw_text):
37
38             #!/*** Findall Non text characters in each word
39             non_text = re.findall("\W", x)
40
41             #!/*** Remove non_text Characters
42             for i in non_text:
43                 x = x.replace(i, "")
44
45             #!/*** If X has length, append out
46             if len(x) > 0:
47                 tokens.append(x.lower())
48         return tokens
49
50     def build_ngrams(self):
51         if self.ngram_size <= 0:
52             print("Ngram size must be an integer > 0")
53             print("Quitting!")
54             return None
55
56         #!/*** Using unigrams, use tokens
57         if self.ngram_size == 1:
58             self.corpus_ngrams = self.corpus_tokens
59             return
60
61         self.corpus_ngrams = []
62
63         #!/*** Get each token group from corpus_tokens
64         for token in self.corpus_tokens:
65
66             loop_ngram = []
67
68             #!/*** Use an index based range to loop through tokens
69             for x in range(0, len(token) ):
70
71                 #!/*** Check if index + ngram_size exceeds the length
```

```
72         if x+self.ngram_size <= len(token):
73
74             result = ""
75
76             #!/*** Build the ngram
77             for y in range(self.ngram_size):
78                 #print(self.tokens[x+y])
79                 result += token[x+y] + " "
80
81             loop_ngram.append(result[:-1])
82
83         else:
84             break
85
86         #!/*** Grab Token Element Count, Keep the greatest count
87         if len(loop_ngram) > self.max_element_count:
88             self.max_element_count = len(loop_ngram)
89
90
91         #!/*** Token group ngram is built. Add loop_ngram to corpus
92         self.corpus_ngrams.append(loop_ngram)
93
94
95
96     def build_vocabulary(self, corpus):
97         if not isinstance(corpus, list):
98             print("Vectorizer Requires a corpus (list of text):")
99             return None
100
101         self.tokens = []
102
103         print("Tokenizing...")
104         #!/*** Tokenize each text entry in the corpus
105         for raw_text in corpus:
106             self.corpus_tokens.append(self.tokenize(raw_text))
107
108         print("Building ngrams...")
109         #!/*** Build ngrams (Defaults to unigrams)
110         self.build_ngrams()
111
112         word_freq = {}
113
114         print("Building Vocabulary...")
115         #!/*** Build dictionary of unique words
116         #!/*** Loop through each element of the corpus
117         for element in self.corpus_ngrams:
118
119
120             #!/*** Grab Token Element Count, Keep the greatest count
121             if len(element) > self.max_element_count:
122                 self.max_element_count = len(element)
123
124             #!/*** Process each individual ngram
125             for ngram in element:
126
127
```

```

128
129         #!/*** Add unique words to dictionaries
130         if ngram not in self.vocabulary_index.keys():
131             index = len(self.ngram_index.values())
132             self.ngram_index[ index ] = ngram
133             self.vocabulary_index [ ngram ] = index
134
135         #!/*** Initialize Word Frequency
136         word_freq[ ngram ] = 1
137     else:
138         #!/*** Increment Word Frequency
139         word_freq[ ngram ] += 1
140
141     #!/*** END for element in self.corpus_ngrams:
142     if self.max_tokens != None:
143
144         #!/*** Check if token count exceeds max tokens
145         if self.max_tokens < len(self.ngram_index.items()):
146
147             print("Sorting Word Frequency...")
148             #!/*** Sort the Word Frequency Dictionary. Keep the 1
149             word_freq = dict(sorted(word_freq.items(), key=lambda
150
151             print("Building Token Dictionary")
152             #!/*** Get list of keys that are lowest frequency
153             for key in list(word_freq.keys())[self.max_tokens:]:
154                 #!/*** Delete Low Frequency ngrams
155                 del word_freq[ key ]
156
157             self.ngram_index = {}
158             self.vocabulary_index = {}
159
160             print("Rebuilding Vocabulary")
161             #!/*** Rebuild ngram_index & vocabulary_index
162             for ngram in word_freq.keys():
163                 index = len(self.ngram_index.values())
164                 self.ngram_index[ index ] = ngram
165                 self.vocabulary_index [ ngram ] = index
166
167             #!/*** END Trim Low Frequency ngrams
168             self.word_freq = word_freq
169
170     #!/*** One Hot encode the corpus.
171     #!/*** Handling the corpus as a whole increases processing speed
172     #!/*** Hot encode to a sparse tensor to for increased encoding sp
173     def one_hot_encode(self, corpus):
174
175         #!/*** Encoded Results
176         results = []
177
178         #!/*** Set the Max array size to the total number of items in
179         array_size = len(self.ngram_index.keys())
180
181
182         start_time = datetime.datetime.now()
183         count = 0

```



```

184
185
186
187     for element in corpus:
188         #!/*** hot encode each ngram
189         result = []
190         for ngram in element:
191
192             #!/*** Skip words not in self.vocabulary_index
193             #!/*** These are skipped due to max_tokens limitation
194             if ngram not in self.vocabulary_index.keys():
195                 continue
196
197             sparse_tensor = tf.SparseTensor(indices=[[0,self.voca
198             #index = self.vocabulary_index[ngram]
199
200             #base_array = np.zeros(array_size, dtype=int)
201
202             #base_array [index] = 1
203
204
205             #!/*** Add the one-hot-encoded word to encoded text
206             result.append(sparse_tensor)
207
208
209             #!/*** END for ngram in tokens:
210
211             result = tf.sparse.concat(axis=1, sp_inputs=result)
212             #!/*** concat Sparse Matrix
213             results.append( result )
214
215             count += 1
216
217
218
219             #!/*** Print a status update every 1000 items
220             if count % 100 == 0:
221                 print(f"{count} / {len(corpus)} Encoded: {datetime.da
222
223             #!/*** Concat List of Sparse Matrixes into a sparse matrix
224             #results = tf.sparse.concat(axis=1, sp_inputs=results)
225
226             print(f"Encoding Complete: {datetime.datetime.now() - start_t
227
228             return results
229
230     #!/***
231     def integer_encode(self,corpus):
232         #!/*** Encoded Results
233         results = []
234
235         #!/*** Set the Max array size to the total number of items in
236         array_size = len(self.ngram_index.keys())
237
238
239         start_time = datetime.datetime.now()

```

```
240         count = 0
241
242
243
244     for element in corpus:
245         #!/*** hot encode each ngram
246         result = []
247         for ngram in element:
248
249             #!/*** Skip words not in self.vocabulary_index
250             #!/*** These are skipped due to max_tokens limitation
251             if ngram not in self.vocabulary_index.keys():
252                 continue
253
254             #!/*** Get integer value of ngram from dictionary.
255             #!/*** Add to result
256             result.append(self.vocabulary_index[ngram])
257
258
259         #!/*** END for ngram in tokens:
260
261         #!/*** result is a complete encoded element
262         results.append( np.array(result).astype(np.float32) )
263
264         count += 1
265
266
267
268         #!/*** Print a status update every 1000 items
269         if count % 5000 == 0:
270             print(f"{count} / {len(corpus)} Encoded: {datetime.da
271
272     print(f"Encoding Complete: {datetime.datetime.now() - start_t
273
274     #!/*** results is a collection of encoded elements
275     return np.array(results, dtype=object)
276
277
278 def encode(self, corpus, encoding='int'):
279
280     if not isinstance(corpus, list) :
281         print("Vectorizer Requires a corpus (list of text):")
282         return None
283
284     self.corpus_tokens = []
285     self.corpus_ngrams = []
286     print("Tokenizing...")
287     #!/*** Tokenize each text entry in the corpus
288     for raw_text in corpus:
289         self.corpus_tokens.append(self.tokenize(raw_text))
290
291     print("Building ngrams...")
292     #!/*** Build ngrams (Defaults to unigrams)
293     self.build_ngrams()
294
295     if encoding == 'onehot':
```

```
296         print("One Hot Coding....")
297
298         #!/*** One Hot Encode Values. These are actually sparse t
299         encoded = self.one_hot_encode(self.corpus_ngrams)
300
301     if encoding == 'int':
302         print("Integer encoding....")
303
304         #!/*** Convert ngrams to integers. These are actually spa
305         encoded = self.integer_encode(self.corpus_ngrams)
306
307         #!/*** Convert lists to Numpy array of float 32 type. Thi
308         #encoded = np.asarray(encoded).astype('float32')
309
310         #!/*** TidyUp (Delete) ngrams and Tokens
311     if self.tidyup:
312         self.corpus_tokens = []
313         self.corpus_ngrams = []
314
315
316     return encoded
317
318     #!/*** Convert One-Hot-Encoding to text
319 def decode(self,elements):
320
321     results = []
322
323     #!/*** For Each element in Corpus
324
325     decoded = ""
326
327     #!/*** For Each ngram (word(s)) in Elements
328     for ngram in elements:
329
330         #!/*** Grab Index of 1 from sparse tensor
331         index = ngram.indices[0].numpy()[1]
332
333         #ngram = list(ngram.numpy())
334
335         decoded += self.ngram_index[ index ] + " "
336
337     #!/*** END for ngram in elements:
338     results.append( decoded[:-1])
339
340     #!/*** END for elements in corpus:
341     return results
342
343 print("Loading Raw Validation Set")
344 val_x_train, val_y_train = load_dataset(val_dir)
345
346 #!/*** Test the Vectorizer with some sample data
347 vectorizer = Vectorizer(max_tokens=100,ngrams=2, tidyup=False)
348 vectorizer.build_vocabulary(val_x_train[:5])
349 start_time = datetime.datetime.now()
350
351 temp_vals = vectorizer.encode(val_x_train[:5],encoding='onehot')
```

```

352
353 print(f"Run Time: {datetime.datetime.now() - start_time}")
354
355
356 int_vals = vectorizer.encode(val_x_train[:5], encoding='int')
357 print("Integer Encoding:")
358 print(int_vals)
359 print()
360 print()
361
362 print("Sample Text: (First 500 Chars)")
363 for element in val_x_train[:5]:
364     print(element[:500])
365     print("====")
366 print()
367 print()
368
369 print("Tokens: (First 100 tokens)")
370 for token in vectorizer.corpus_tokens:
371     print(token[:100])
372     print("====")
373 print()
374 print()
375
376 print("ngrams: (First 50 tokens)")
377 for token in vectorizer.corpus_ngrams:
378     print(token[:100])
379     print("====")
380 print()
381 print()
382 print("Small one hot encoded Sample:")
383 print(temp_vals)
384 print()
385 print()
386 print("Encoded Vocabulary")
387 print(vectorizer.vocabulary_index)
388 print()
389 print()
390 print("Decoded Text from vocabulary (limited by max tokens)")
391 for result in vectorizer.decode(temp_vals):
392     print(result)
393     print()
394
395 del temp_vals
Loading Raw Validation Set
396 del int_vals
Tokenizing...
397 del vectorizer
Building ngrams...
Building Vocabulary...
Sorting Word Frequency...
Building Token Dictionary
Rebuilding Vocabulary
Tokenizing...
Building ngrams...
One Hot Coding....
Encoding Complete: 0:00:00.046011
Run Time: 0:00:00.051012
Tokenizing...

```

```

Building ngrams...
Integer encoding....
Encoding Complete: 0:00:00
Integer Encoding:
[array([ 2., 76., 77.,  5., 78., 79., 80., 81., 82., 21., 83., 84., 8
5.,

```

```

In [7]: 1 print("Loading Raw Validation Set")
        2 raw_val_x_train, val_y_train = load_dataset(val_dir)
        3
        4 print("Loading Raw Train Data")
        5 raw_x_train, y_train = load_dataset(train_dir)
        6
        7 print("Loading Raw Test Data")
        8 raw_x_test, y_test = load_dataset(test_dir)
        9 print("Done")
       10
       11 val_y_train = np.array(val_y_train, dtype=object)
       12 y_train = np.array(y_train, dtype=object)
       13 y_test = np.array(y_test, dtype=object)
       14
       15 y_train = np.asarray(y_train).astype(np.int)
       16

```

```

Loading Raw Validation Set
Loading Raw Train Data
Dropping File: 7714_1.txt due to decoding issues
Dropping File: 11351_9.txt due to decoding issues
Dropping File: 8263_9.txt due to decoding issues
Loading Raw Test Data
Dropping File: 4414_1.txt due to decoding issues
Dropping File: 6973_1.txt due to decoding issues
Dropping File: 2464_10.txt due to decoding issues
Dropping File: 5281_10.txt due to decoding issues
Done

```

```

In [8]: 1
        2
        3
        4 #!/*** Test the Vectorizer with some sample data
        5 max_tokens = 20000
        6 #maxlen = 1000 #!/*** Limit reviews to this length, Leave blank to us
        7
        8 ngrams = 1
        9
       10 #!/*** Initialize vectorizer
       11 vectorizer = Vectorizer(max_tokens=max_tokens, ngrams=ngrams)
       12
       13 #!/*** Build Vocabulary based on the training text
       14 vectorizer.build_vocabulary(raw_x_train)
       15
       16 #!/*** maxlen required: This is maximum number of tokens/ngrams to us
       17 #!/*** pads all articles to the same word count. This required to hav
       18
       19 maxlen = vectorizer.max_element_count #!/*** Sets maxlen to the large
       20
       21

```

```

22 #!/*** Encode Validation, training and test data
23
24 print("Encoding Validation Data...")
25 val_x_train = vectorizer.encode(raw_val_x_train)
26 print(vectorizer.max_element_count)
27 print("Encoding Training Data...")
28 x_train = vectorizer.encode(raw_x_train)
29
30 print("Encoding Test Data...")
31 x_test = vectorizer.encode(raw_x_test)
32
33
34 print("Padding Validation...")
35 val_x_train = preprocessing.sequence.pad_sequences(val_x_train, maxlen=
36
37 print("Padding Training Data...")
38 x_train = preprocessing.sequence.pad_sequences(x_train, maxlen=maxlen
39
40 print("Padding Test Data...")
41 x_test = preprocessing.sequence.pad_sequences(x_test, maxlen=maxlen)
42
43 val_y_train = np.asarray(val_y_train).astype(np.int)
44 x_test = np.asarray(x_test).astype(np.int)
45 y_test = np.asarray(y_test).astype(np.int)
46
47 print("Done")
48
49 print(type(val_x_train),type(val_x_train[0]),x_train.shape)
50 print(type(val_y_train),type(val_y_train[0]),y_train.shape)
51 print(type(x_test),type(x_test[0]))
52 print(type(y_test),type(y_test[0]))
53 print(x_test.shape,y_test.shape)
54

```

```

Tokenizing...
Building ngrams...
Building Vocabulary...
Sorting Word Frequency...
Building Token Dictionary
Rebuilding Vocabulary
Encoding Validation Data...
Tokenizing...
Building ngrams...
Integer encoding....
5000 / 5000 Encoded: 0:00:00.496111
Encoding Complete: 0:00:00.496111
2450
Encoding Training Data...
Tokenizing...
Building ngrams...
Integer encoding....
5000 / 19997 Encoded: 0:00:00.484107
10000 / 19997 Encoded: 0:00:00.977218
15000 / 19997 Encoded: 0:00:01.466329
Encoding Complete: 0:00:01.963440
Encoding Test Data...
Tokenizing...

```

```

Building ngrams...
Integer encoding....
5000 / 24996 Encoded: 0:00:00.477108
10000 / 24996 Encoded: 0:00:00.949213
15000 / 24996 Encoded: 0:00:01.431323
20000 / 24996 Encoded: 0:00:01.913430
Encoding Complete: 0:00:02.384537
Padding Validation...
Padding Training Data...
Padding Test Data...
Done
<class 'numpy.ndarray'> <class 'numpy.ndarray'> (19997, 2450)
<class 'numpy.ndarray'> <class 'numpy.int32'> (19997,)
<class 'numpy.ndarray'> <class 'numpy.ndarray'>
<class 'numpy.ndarray'> <class 'numpy.int32'>
(24996, 2450) (24996,)

```

```

In [9]: 1 print(type(val_x_train),type(val_x_train[0]),x_train.shape)
        2 print(type(val_y_train),type(val_y_train[0]),y_train.shape)
        3 print(type(x_test),type(x_test[0]))
        4 print(type(y_test),type(y_test[0]))
        5 print(x_test.shape,y_test.shape)

```

```

<class 'numpy.ndarray'> <class 'numpy.ndarray'> (19997, 2450)
<class 'numpy.ndarray'> <class 'numpy.int32'> (19997,)
<class 'numpy.ndarray'> <class 'numpy.ndarray'>
<class 'numpy.ndarray'> <class 'numpy.int32'>
(24996, 2450) (24996,)

```

https://github.com/fchollet/deep-learning-with-python-notebooks/blob/master/first_edition/6.1-using-word-embeddings.ipynb (https://github.com/fchollet/deep-learning-with-python-notebooks/blob/master/first_edition/6.1-using-word-embeddings.ipynb)

In [10]:

```

1  """
2
3  from keras.datasets import imdb
4  from keras import preprocessing
5
6  # Number of words to consider as features
7  max_features = 10000
8  # Cut texts after this number of words
9  # (among top max_features most common words)
10 maxlen = 20
11
12 # Load the data as lists of integers.
13 (x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=max_f
14
15 # This turns our lists of integers
16 # into a 2D integer tensor of shape `(samples, maxlen)`
17 #x_train = preprocessing.sequence.pad_sequences(x_train, maxlen=maxle
18 #x_test = preprocessing.sequence.pad_sequences(x_test, maxlen=maxlen)
19 print(type(x_train), type(x_train[0]))
20 print(type(y_train), type(y_train[0]))
21 """
22 print()

```

pad_sequence: Used to transform lists within lists (2D array) to have a uniform inner dimension. Essentially, padding smaller arrays to the size of the largest, or trimming all arrays to max len

This function transforms a list (of length num_samples) of sequences (lists of integers) into a 2D Numpy array of shape (num_samples, num_timesteps). num_timesteps is either the maxlen argument if provided, or the length of the longest sequence in the list.

- https://www.tensorflow.org/api_docs/python/tf/keras/preprocessing/sequence/pad_sequences (https://www.tensorflow.org/api_docs/python/tf/keras/preprocessing/sequence/pad_sequences)

10.2

Using listings 6.16, 6.17, and 6.18 in Deep Learning with Python as a guide, train a sequential model with embeddings on the IMDB data found in data/external/imdb/. Produce the model performance metrics and training and validation accuracy curves within the Jupyter notebook.

In [11]:

```

1
2
3  def get_sequential_embedded_model():
4      model = Sequential()
5      # We specify the maximum input length to our Embedding layer
6      # so we can later flatten the embedded inputs
7      model.add(Embedding(max_tokens, 8, input_length=maxlen))
8      # After the Embedding layer,
9      # our activations have shape `(samples, maxlen, 8)`.
10

```



```

11     # We flatten the 3D tensor of embeddings
12     # into a 2D tensor of shape `(samples, maxlen * 8)`
13     model.add(Flatten())
14
15     # We add the classifier on top
16     model.add(Dense(1, activation='sigmoid'))
17     model.compile(optimizer='rmsprop', loss='binary_crossentropy', me
18
19     return model
20
21 model = get_sequential_embedded_model()
22 model.summary()
23
24 #!/*** Code to check the data type expected for each model layer
25 [print(i.shape, i.dtype) for i in model.inputs]
26 [print(o.shape, o.dtype) for o in model.outputs]
27 [print(l.name, l.input_shape, l.dtype) for l in model.layers]
28
29 #!/*** Reference for recasting lists as an np.array of float32 type
30 #x = np.asarray(x).astype('float32')
31
32

```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
embedding (Embedding)	(None, 2450, 8)	160000

flatten (Flatten)	(None, 19600)	0

dense (Dense)	(None, 1)	19601
=====		

Total params: 179,601
Trainable params: 179,601
Non-trainable params: 0

```

(None, 2450) <dtype: 'float32'>
(None, 1) <dtype: 'float32'>
embedding (None, 2450) float32
flatten (None, 2450, 8) float32
dense (None, 19600) float32

```

Out[11]: [None, None, None]

```



In [12]: 1 history = model.fit(val_x_train, val_y_train,
2                               epochs=10,
3                               batch_size=32,
4                               validation_split=0.2
5                               )
6
7 score = model.evaluate(np.asarray(x_test).astype(np.int), np.asarray(y_test).astype(np.int))
8
9 print('Test loss:', score[0])
10 print('Test accuracy:', score[1])
11
12 plot_model_history(history)


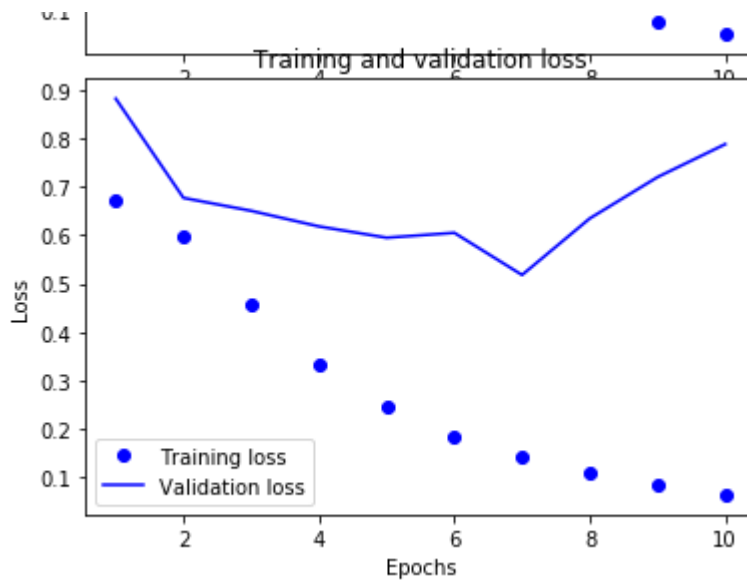
```

```

Epoch 1/10
125/125 [=====] - 2s 9ms/step - loss: 0.6708
- acc: 0.6208 - val_loss: 0.8819 - val_acc: 0.0000e+00
Epoch 2/10
125/125 [=====] - 1s 7ms/step - loss: 0.5973
- acc: 0.6758 - val_loss: 0.6769 - val_acc: 0.5780
Epoch 3/10
125/125 [=====] - 1s 8ms/step - loss: 0.4582
- acc: 0.8303 - val_loss: 0.6502 - val_acc: 0.6390
Epoch 4/10
125/125 [=====] - 1s 7ms/step - loss: 0.3333
- acc: 0.8988 - val_loss: 0.6180 - val_acc: 0.6700
Epoch 5/10
125/125 [=====] - 1s 7ms/step - loss: 0.2458
- acc: 0.9335 - val_loss: 0.5947 - val_acc: 0.6950
Epoch 6/10
125/125 [=====] - 1s 7ms/step - loss: 0.1850
- acc: 0.9523 - val_loss: 0.6048 - val_acc: 0.6910
Epoch 7/10
125/125 [=====] - 1s 7ms/step - loss: 0.1410
- acc: 0.9685 - val_loss: 0.5178 - val_acc: 0.7540
Epoch 8/10
125/125 [=====] - 1s 7ms/step - loss: 0.1088
- acc: 0.9760 - val_loss: 0.6345 - val_acc: 0.7030
Epoch 9/10
125/125 [=====] - 1s 7ms/step - loss: 0.0836
- acc: 0.9822 - val_loss: 0.7198 - val_acc: 0.6830
Epoch 10/10
125/125 [=====] - 1s 7ms/step - loss: 0.0642
- acc: 0.9883 - val_loss: 0.7878 - val_acc: 0.6700
Test loss: 0.4089915454387665
Test accuracy: 0.9270724717140100

```





```
In [13]: 1 history = model.fit(x_train,y_train,
2                       epochs=8,
3                       batch_size=32,
4                       validation_split=0.2
5                       )
6
7 score = model.evaluate(np.asarray(x_test).astype(np.int), np.asarray(y_test).astype(np.int))
8
9 print('Test loss:', score[0])
10 print('Test accuracy:', score[1])
11
12
```

```
Epoch 1/8
500/500 [=====] - 4s 8ms/step - loss: 0.2878
- acc: 0.8827 - val_loss: 0.4963 - val_acc: 0.7975
Epoch 2/8
500/500 [=====] - 4s 8ms/step - loss: 0.2160
- acc: 0.9155 - val_loss: 0.5547 - val_acc: 0.7803
Epoch 3/8
500/500 [=====] - 4s 8ms/step - loss: 0.1730
- acc: 0.9362 - val_loss: 0.4601 - val_acc: 0.8198
Epoch 4/8
500/500 [=====] - 4s 8ms/step - loss: 0.1423
- acc: 0.9492 - val_loss: 0.4608 - val_acc: 0.8290
Epoch 5/8
500/500 [=====] - 4s 8ms/step - loss: 0.1186
- acc: 0.9595 - val_loss: 0.6408 - val_acc: 0.7710
Epoch 6/8
500/500 [=====] - 4s 8ms/step - loss: 0.0997
```

In []:

10.3

Using listing 6.27 in Deep Learning with Python as a guide, fit the same data with an LSTM layer. Produce the model performance metrics and training and validation accuracy curves within the Jupyter notebook.

In [14]:

```

1
2
3 from keras.models import Sequential
4 from keras.layers import Embedding, SimpleRNN, LSTM
5
6 def get_lstm_model():
7
8     model = Sequential()
9     model.add(Embedding(max_tokens, 32))
10    model.add(LSTM(32))
11    model.add(Dense(1, activation='sigmoid'))
12
13    model.compile(optimizer='rmsprop',
14                  loss='binary_crossentropy',
15                  metrics=['acc'])
16
17
18 history = model.fit(val_x_train, val_y_train,
19                     epochs=10,
20                     batch_size=128,
21                     validation_split=0.2)
22
23 score = model.evaluate(np.asarray(x_test).astype(np.int), np.asarray(y_test).astype(np.int))
24
25 print('Test loss:', score[0])
26 print('Test accuracy:', score[1])
27
28 plot_model_history(history)

```

Epoch 1/10

```

32/32 [=====] - 1s 27ms/step - loss: 0.0881 -
acc: 0.9685 - val_loss: 0.7499 - val_acc: 0.7710

```

Epoch 2/10

```

32/32 [=====] - 1s 26ms/step - loss: 0.0703 -
acc: 0.9770 - val_loss: 0.7489 - val_acc: 0.7780

```

Epoch 3/10

```

32/32 [=====] - 1s 26ms/step - loss: 0.0569 -
acc: 0.9830 - val_loss: 0.8917 - val_acc: 0.7380

```

Epoch 4/10

```

32/32 [=====] - 1s 26ms/step - loss: 0.0464 -
acc: 0.9875 - val_loss: 0.8226 - val_acc: 0.7510

```

Epoch 5/10

```

32/32 [=====] - 1s 26ms/step - loss: 0.0377 -
acc: 0.9902 - val_loss: 0.8585 - val_acc: 0.7510

```

Epoch 6/10

```

32/32 [=====] - 1s 26ms/step - loss: 0.0313 -
acc: 0.9933 - val_loss: 1.0313 - val_acc: 0.7050

```

Epoch 7/10

```

32/32 [=====] - 1s 27ms/step - loss: 0.0255 -
acc: 0.9950 - val_loss: 0.9157 - val_acc: 0.7430

```

Epoch 8/10

```

32/32 [=====] - 1s 26ms/step - loss: 0.0208 -
acc: 0.9965 - val_loss: 0.9251 - val_acc: 0.7420

```

Epoch 9/10

```

32/32 [=====] - 1s 26ms/step - loss: 0.0171 -
acc: 0.9975 - val_loss: 1.0651 - val_acc: 0.7080

```



Epoch 10/10


32/32 [=====] - 1s 26ms/step - loss: 0.0135 -

acc: 0.9977 - val_loss: 1.0511 - val_acc: 0.7150

Test loss: 0.4457803964614868

Test accuracy: 0.8521264708545827





```
In [15]: 1 history = model.fit(x_train, y_train,
2                       epochs=8,
3                       batch_size=128,
4                       validation_split=0.2)
5
6 score = model.evaluate(np.asarray(x_test).astype(np.int), np.asarray(y_test).astype(np.int))
7
8 print('Test loss:', score[0])
9 print('Test accuracy:', score[1])
```

```
Epoch 1/8
125/125 [=====] - 3s 27ms/step - loss: 0.0648
- acc: 0.9773 - val_loss: 0.8485 - val_acc: 0.7442
Epoch 2/8
125/125 [=====] - 3s 26ms/step - loss: 0.0523
- acc: 0.9834 - val_loss: 0.7618 - val_acc: 0.7663
Epoch 3/8
125/125 [=====] - 3s 26ms/step - loss: 0.0432
- acc: 0.9874 - val_loss: 0.8154 - val_acc: 0.7508
Epoch 4/8
125/125 [=====] - 3s 26ms/step - loss: 0.0359
- acc: 0.9905 - val_loss: 0.7709 - val_acc: 0.7663
Epoch 5/8
125/125 [=====] - 3s 26ms/step - loss: 0.0297
- acc: 0.9924 - val_loss: 0.8048 - val_acc: 0.7607
Epoch 6/8
125/125 [=====] - 3s 26ms/step - loss: 0.0247
- acc: 0.9942 - val_loss: 0.9268 - val_acc: 0.7372
Epoch 7/8
125/125 [=====] - 3s 26ms/step - loss: 0.0203
- acc: 0.9955 - val_loss: 0.9174 - val_acc: 0.7420
Epoch 8/8
125/125 [=====] - 3s 26ms/step - loss: 0.0167
- acc: 0.9966 - val_loss: 0.8256 - val_acc: 0.7688
Test loss: 0.46850863099098206
Test accuracy: 0.8548167943954468
```

10.4

Using listing 6.46 in Deep Learning with Python as a guide, fit the same data with a simple 1D convnet. Produce the model performance metrics and training and validation accuracy curves within the Jupyter notebook.

```
In [24]: 1 from keras import layers
2
3 def get_conv1d_model():
4     model = Sequential()
5     model.add(layers.Embedding(max_tokens, 128, input_length=maxlen))
6     model.add(layers.Conv1D(32, 7, activation='relu'))
7     model.add(layers.MaxPooling1D(5))
8     model.add(layers.Conv1D(32, 7, activation='relu'))
9     model.add(layers.GlobalMaxPooling1D())
10    model.add(layers.Dense(1))
11
12    model.compile(optimizer=RMSprop(learning_rate=1e-4),
13                  loss='binary_crossentropy',
14                  metrics=['acc'])
15
16    model.summary()
17    return model
18
19
20 model = get_conv1d_model()
21 history = model.fit(val_x_train, val_y_train,
22                     epochs=30,
23                     batch_size=128,
24                     validation_split=0.2)
25
26 score = model.evaluate(np.asarray(x_test).astype(np.int), np.asarray(y_test))
27
28 print('Test loss:', score[0])
29 print('Test accuracy:', score[1])
30
31 plot_model_history(history)
```


Model: "sequential_6"

Layer (type)	Output Shape	Param #
--------------	--------------	---------

```
In [25]: 1 history = model.fit(x_train, y_train,
2                             epochs=10,
3                             batch_size=128,
4                             validation_split=0.2)
5
6 score = model.evaluate(np.asarray(x_test).astype(np.int), np.asarray(y_test).astype(np.int))
7
8 print('Test loss:', score[0])
9 print('Test accuracy:', score[1])
10
```

Epoch 1/10

125/125 [=====] - 99s 789ms/step - loss: 3.1318 - acc: 0.6251 - val_loss: 1.2540 - val_acc: 0.0000e+00

Epoch 2/10

125/125 [=====] - 97s 775ms/step - loss: 0.6554 - acc: 0.6251 - val_loss: 0.9645 - val_acc: 0.0000e+00

Epoch 3/10

125/125 [=====] - 99s 789ms/step - loss: 0.6306 - acc: 0.6251 - val_loss: 0.9557 - val_acc: 2.5000e-04

Epoch 4/10

125/125 [=====] - 101s 809ms/step - loss: 0.5968 - acc: 0.6288 - val_loss: 0.9135 - val_acc: 0.0207

Epoch 5/10

125/125 [=====] - 98s 786ms/step - loss: 0.5406 - acc: 0.6777 - val_loss: 0.9148 - val_acc: 0.1615

Epoch 6/10

125/125 [=====] - 99s 794ms/step - loss: 0.4434 - acc: 0.8117 - val_loss: 0.8259 - val_acc: 0.4803

Epoch 7/10

125/125 [=====] - 98s 781ms/step - loss: 0.3456 - acc: 0.8776 - val_loss: 0.4988 - val_acc: 0.7897

Epoch 8/10

125/125 [=====] - 97s 779ms/step - loss: 0.2825 - acc: 0.9052 - val_loss: 0.6948 - val_acc: 0.7120

Epoch 9/10

125/125 [=====] - 97s 779ms/step - loss: 0.2384 - acc: 0.9191 - val_loss: 0.7116 - val_acc: 0.7322

Epoch 10/10




125/125 [=====] - 100s 798ms/step - loss: 0.2044 - acc: 0.9331 - val_loss: 0.7217 - val_acc: 0.7487

Test loss: 0.4732949435710907

Test accuracy: 0.8379340767860413

In [26]:

```
1 plot_model.history(history)
```



Training and validation accuracy

