

Stoneburner, Kurt

. DSC 650 - Tensorflow Keras Binary Classifier Example

```
In [1]: 1 import os
2 import sys
3 # Imports and Load Data
4 import matplotlib.pyplot as plt
5 import numpy as np
6 import pandas as pd
7 from tensorflow.keras.datasets import imdb
8 from tensorflow.keras import models
9 from tensorflow.keras import layers
10 from tensorflow.keras import optimizers
11
12
13 Use the whole window in the IPYNB editor
14 from IPython.core.display import display, HTML
15 display(HTML("<style>.container { width:100% !important; }</style>"))
16
17 Maximize columns and rows displayed by pandas
18 pd.set_option('display.max_rows', 100)
```

```
In [2]: 1 from os import environ
2
3 environ["KERAS_BACKEND"] = "plaidml.keras.backend"
4
```

```
In [3]: 1 Download Data and Load arrays
```

<__array_function__ internals>:5: VisibleDeprecationWarning: Creating an ndarray from ragged nested sequences (which is a list-or-tuple of lists-or-tuples-or ndarrays with different lengths or shapes) is deprecated. If you meant to do this, you must specify 'dtype=object' when creating the ndarray.

C:\Users\stonk013\Anaconda3\envs\keras_env\lib\site-packages\tensorflow\python\keras\datasets\imdb.py:159: VisibleDeprecationWarning: Creating an ndarray from ragged nested sequences (which is a list-or-tuple of lists-or-tuples-or ndarrays with different lengths or shapes) is deprecated. If you meant to do this, you must specify 'dtype=object' when creating the ndarray.

```
    x_train, y_train = np.array(xs[:idx]), np.array(labels[:idx])
```

C:\Users\stonk013\Anaconda3\envs\keras_env\lib\site-packages\tensorflow\python\keras\datasets\imdb.py:160: VisibleDeprecationWarning: Creating an ndarray from ragged nested sequences (which is a list-or-tuple of lists-or-tuples-or ndarrays with different lengths or shapes) is deprecated. If you meant to do this, you must specify 'dtype=object' when creating the ndarray.

```
    x_test, y_test = np.array(xs[idx:]), np.array(labels[idx:])
```

In []: ▶

In [4]: ▶

```

1
2 #!/*** Word_index is dictionary mapping words to an integer index
3 word_index = imdb.get_word_index()
4
5 #!/*** Maps indexes to words
6 reverse_word_index = dict(
7     [(value, key) for (key, value) in word_index.items()])
8
9 #!/*** Decodes the Review
10 decoded_review = ' '.join(

```

In []: ▶

You can't feed lists of integers into a neural network. You have to turn your lists into tensors. There are two ways to do that:

- Pad your lists so that they all have the same length, turn them into an integer tensor of shape (samples, word_indices), and then use as the first layer in your network a layer capable of handling such integer tensors (the Embedding layer, which we'll cover in detail later in the book).
- **(Example Below: vectorize_sequences)** One-hot encode your lists to turn them into vectors of 0s and 1s. This would mean, for instance, turning the sequence [3, 5] into a 10,000-dimensional vector that would be all 0s except for indices 3 and 5, which would be 1s. Then you could use as the first layer in your network a Dense layer, capable of handling floating-point vector data.

In [5]: ▶

```

1 #!/*** Lists of integers must be converted into tensors.
2 def vectorize_sequences(sequences, dimension=10000):
3     #!/*** Builds zero filled matrix of shape dimension
4     results = np.zeros((len(sequences), dimension))
5
6     #!/*** Assigns 1s to the specific integer for references.
7     #!/*** This is manual one-hot encoding
8     for i, sequence in enumerate(sequences):
9
10         for j in sequence:
11             results[i, j] = 1.
12
13     return results
14
15 #!/*** Download Data and Load arrays
16 (train_data, train_labels), (test_data, test_labels) = imdb.load_data
17
18 #!/*** Vectorize the Training and Test data
19 x_train = vectorize_sequences(train_data)
20 x_test = vectorize_sequences(test_data)
21
22 #!/*** Vectorize the Training and Test Labels

```

```
23 y_train = np.asarray(train_labels).astype('float32')
```

Build the Network

The input data is vectors, and the labels are scalars (1s and 0s): this is the easiest setup you'll ever encounter. A type of network that performs well on such a problem is a simple stack of fully connected (Dense) layers with relu activations: *Dense(16,activation='relu')*.

The argument being passed to each Dense layer (16) is the number of hidden units of the layer. A hidden unit is a dimension in the representation space of the layer. Each such Dense layer with a relu activation implements the following chain of tensor operations:

output = relu(dot(W, input) + b)

Having 16 hidden units means the weight matrix W will have shape (input_dimension,16): the dot product with W will project the input data onto a 16-dimensional representation space (and then you'll add the bias vector b and apply the relu operation).

You can intuitively understand the dimensionality of your representation space as “how much freedom you’re allowing the network to have when learning internal representations.” *Having more hidden units (a higher-dimensional representation space) allows your network to learn more-complex representations, but it makes the network more computationally expensive and may lead to learning unwanted patterns (patterns that will improve performance on the training data but not on the test data).*

There are two key architecture decisions to be made about such a stack of Dense layers:

- How many layers to use
- How many hidden units to choose for each layer

For the time being go with the following architecture choice:

- Two intermediate layers with 16 hidden units each
- A third layer that will output the scalar prediction regarding the sentiment of the current review

The intermediate layers will use *relu* as their activation function, and the final layer will use a sigmoid activation so as to output a probability (a score between 0 and 1, indicating how likely the sample is to have the target “1”: how likely the review is to be positive). A *relu* (rectified linear unit) is a function meant to zero out negative values, whereas a sigmoid “squashes” arbitrary values into the [0, 1] interval, outputting something that can be interpreted as a probability

```
In [6]: ▶ 1
          2 model = models.Sequential()
          3 model.add(layers.Dense(16, activation='relu', input_shape=(10000,))
          4 model.add(layers.Dense(16, activation='relu'))
          5 model.add(layers.Dense(1, activation='sigmoid'))
          6
          7 #!/** Compile the Model
          8
          9
         10
         11 model.compile(
```

```
12     optimizer=optimizers.RMSprop(lr=0.001),
13     loss='binary_crossentropy',
14     metrics=['accuracy']
15 )
16
17 #!/*** Configure the Optimizer
18 #model.compile(
19 #     optimizer=optimizers.RMSprop(lr=0.001),
20 #     loss=losses.binary_crossentropy,
21 #     metrics=[metrics.binary_accuracy]
22 #)
23
24
25 #!/*** Use Custom Losses and Metrics
26 from tensorflow.keras import losses
27 from tensorflow.keras import metrics
28
29 #!/*** Set Aside a Validation Set
30 x_val = x_train[:10000]
31 partial_x_train = x_train[10000:]
32
33 y_val = y_train[:10000]
34 partial_y_train = y_train[10000:]
35
36 #!/*** Train Model
37 model.compile(
38     optimizer='rmsprop',
39     loss='binary_crossentropy',
40     metrics=['acc']
41 )
42
43 history = model.fit(
44     partial_x_train,
45     partial_y_train,
46     epochs=20,
47     batch_size=512,
```

Epoch 1/20

30/30 [=====] - 1s 22ms/step - loss: 0.5176
- acc: 0.7755 - val_loss: 0.4462 - val_acc: 0.7928

Epoch 2/20

30/30 [=====] - 0s 10ms/step - loss: 0.3044
- acc: 0.9021 - val_loss: 0.3056 - val_acc: 0.8882

Epoch 3/20

30/30 [=====] - 0s 10ms/step - loss: 0.2242
- acc: 0.9257 - val_loss: 0.2771 - val_acc: 0.8926

Epoch 4/20

30/30 [=====] - 0s 11ms/step - loss: 0.1767
- acc: 0.9417 - val_loss: 0.2825 - val_acc: 0.8852

Epoch 5/20

30/30 [=====] - 0s 10ms/step - loss: 0.1452
- acc: 0.9531 - val_loss: 0.2785 - val_acc: 0.8908

Epoch 6/20


30/30 [=====] - 0s 10ms/step - loss: 0.1202
- acc: 0.9635 - val_loss: 0.2968 - val_acc: 0.8861

Epoch 7/20

30/30 [=====] - 0s 10ms/step - loss: 0.1008


```
- acc: 0.9705 - val_loss: 0.3074 - val_acc: 0.8827
Epoch 8/20
30/30 [=====] - 0s 10ms/step - loss: 0.0836
- acc: 0.9768 - val_loss: 0.3281 - val_acc: 0.8821
Epoch 9/20
30/30 [=====] - 0s 10ms/step - loss: 0.0714
- acc: 0.9789 - val_loss: 0.3535 - val_acc: 0.8776
Epoch 10/20
30/30 [=====] - 0s 10ms/step - loss: 0.0590
- acc: 0.9844 - val_loss: 0.3713 - val_acc: 0.8790
Epoch 11/20
30/30 [=====] - 0s 10ms/step - loss: 0.0482
- acc: 0.9883 - val_loss: 0.3954 - val_acc: 0.8757
Epoch 12/20
30/30 [=====] - 0s 10ms/step - loss: 0.0441
- acc: 0.9897 - val_loss: 0.4227 - val_acc: 0.8734
Epoch 13/20
30/30 [=====] - 0s 10ms/step - loss: 0.0341
- acc: 0.9930 - val_loss: 0.4499 - val_acc: 0.8718
Epoch 14/20
30/30 [=====] - 0s 10ms/step - loss: 0.0244
- acc: 0.9963 - val_loss: 0.5105 - val_acc: 0.8634
Epoch 15/20
30/30 [=====] - 0s 11ms/step - loss: 0.0233
- acc: 0.9956 - val_loss: 0.5066 - val_acc: 0.8693
Epoch 16/20
30/30 [=====] - 0s 10ms/step - loss: 0.0155
- acc: 0.9984 - val_loss: 0.5628 - val_acc: 0.8631
Epoch 17/20
30/30 [=====] - 0s 10ms/step - loss: 0.0145
- acc: 0.9983 - val_loss: 0.5830 - val_acc: 0.8660
Epoch 18/20
30/30 [=====] - 0s 10ms/step - loss: 0.0138
- acc: 0.9971 - val_loss: 0.6116 - val_acc: 0.8660
Epoch 19/20
30/30 [=====] - 0s 11ms/step - loss: 0.0064
- acc: 0.9996 - val_loss: 0.6432 - val_acc: 0.8650
Epoch 20/20
30/30 [=====] - 0s 10ms/step - loss: 0.0089
- acc: 0.9985 - val_loss: 0.6731 - val_acc: 0.8653
```

```
In [12]: ▶ 1 #!/*** Plotting the training and validation loss
2 history_dict = history.history
3 loss_values = history_dict['loss']
4 val_loss_values = history_dict['val_loss']
5 acc = history_dict['acc']
6 epochs = range(1, len(acc) + 1)
7 plt.plot(epochs, loss_values, 'bo', label='Training loss')
8 plt.plot(epochs, val_loss_values, 'b', label='Validation loss')
9 plt.title('Training and validation loss')
10 plt.xlabel('Epochs')
11 plt.ylabel('Loss')
12 plt.legend()
13 plt.show()
```



```
In [ ]: ▶
```

```
In [13]: ▶ 1 plt.clf()
2 acc_values = history_dict['acc']
3 val_acc_values = history_dict['val_acc']
4 plt.plot(epochs, acc_values, 'bo', label='Training acc')
5 plt.plot(epochs, val_acc_values, 'b', label='Validation acc')
6 plt.title('Training and validation accuracy')
7 plt.xlabel('Epochs')
8 plt.ylabel('Loss')
9 plt.legend()
```



```
In [14]: ▶ 1 #!/*** retraining a model from scratch
2 #!/*** Previous Results didn't go well. Book suggests Overfitting,
3 #!/*** Accuracy and Validation values worked well.
4
5 model = models.Sequential()
6 model.add(layers.Dense(16, activation='relu', input_shape=(10000,)))
7 model.add(layers.Dense(16, activation='relu'))
8 model.add(layers.Dense(1, activation='sigmoid'))
9 model.compile(optimizer='rmsprop', loss='binary_crossentropy', metri
10 model.fit(x_train, y_train, epochs=20, batch_size=512)
11 results = model.evaluate(x_test, y_test)
12
```

```

Epoch 1/20
49/49 [=====] - 3s 15ms/step - loss: 0.4363
- accuracy: 0.8263
Epoch 2/20
49/49 [=====] - 1s 10ms/step - loss: 0.2505
- accuracy: 0.9109
Epoch 3/20
49/49 [=====] - 1s 11ms/step - loss: 0.1930
- accuracy: 0.9325
Epoch 4/20
49/49 [=====] - 1s 12ms/step - loss: 0.1635
- accuracy: 0.9431
Epoch 5/20
49/49 [=====] - 0s 9ms/step - loss: 0.1384 -
accuracy: 0.9535
Epoch 6/20
49/49 [=====] - 1s 11ms/step - loss: 0.1228
- accuracy: 0.9582 0s - loss: 0.1166 - ac
Epoch 7/20
49/49 [=====] - 0s 9ms/step - loss: 0.1100 -
accuracy: 0.9635
Epoch 8/20
49/49 [=====] - 0s 9ms/step - loss: 0.0953 -
accuracy: 0.9691
Epoch 9/20
49/49 [=====] - 0s 9ms/step - loss: 0.0816 -
accuracy: 0.9749
Epoch 10/20
49/49 [=====] - 0s 8ms/step - loss: 0.0747 -
accuracy: 0.9764
Epoch 11/20
49/49 [=====] - 0s 8ms/step - loss: 0.0626 -
accuracy: 0.9814
Epoch 12/20
49/49 [=====] - 0s 8ms/step - loss: 0.0569 -
accuracy: 0.9824
Epoch 13/20
49/49 [=====] - 0s 9ms/step - loss: 0.0506 -
accuracy: 0.9847
Epoch 14/20
49/49 [=====] - 1s 11ms/step - loss: 0.0391
- accuracy: 0.9899
Epoch 15/20
49/49 [=====] - 0s 9ms/step - loss: 0.0352 -
accuracy: 0.9905
Epoch 16/20
49/49 [=====] - 0s 9ms/step - loss: 0.0301 -

```

Out[14]: [0.7695156931877136, 0.8503599762916565]

```

In [17]: ▶ 1 def build_binary_output_model(**kwargs):
2           #!/*** Define the Model
3           from tensorflow.keras import models
4           from tensorflow.keras import layers
5           from tensorflow.keras import optimizers
6
7

```



```
8      #####
9      Set Default values
10     #####
11     total_layers = 2
12     hidden_units = 16
13     first_activation = "relu"
14     final_activation='sigmoid'
15     optimizer='rmsprop'
16     loss = 'mse'
17     metrics=['accuracy']
18     shape = (0,0)
19     do_compile = True
20
21     Apply Kwargs
22     for key,value in kwargs.items():
23
24         if key == 'layers':
25             total_layers=value
26
27         if key == 'hidden_units':
28             hidden_units=value
29
30         if key == 'loss':
31             loss=value
32
33         if key == 'first_activation':
34             first_activation=value
35
36         if key == 'final_activation':
37             final_activation=value
38
39         if key == 'optimizer':
40             optimizer=value
41
42         if key == 'metrics':
43             metrics=value
44
45         if key == 'shape':
46             shape = value
47
48         if key == 'compile':
49             do_compile = value
50
51     model = models.Sequential()
52
53
54     Add First Layer
55     model.add(layers.Dense(hidden_units, activation=first_activation))
56
57
58     Add Additional Layers if total_layers greater than 2
59     for x in range(total_layers-2):
60
61         These are basic layers with same number of hidden units
62         model.add(layers.Dense(hidden_units, activation=first_activation))
63
```

```
64
65
66     #!/*** Add Final Layer
67     model.add(layers.Dense(1, activation=final_activation))
68
69     #!/*** Compile Model
70     if do_compile:
71
72         model.compile(optimizer=optimizer, loss=loss, metrics=metric
73
74     return model
75
```

In [19]:

```
1  #!/*****
2  #!/*** Book Supplied Settings
3  #!/*****
4  model_shape = (10000,)
5  layers = 3
6  hidden_units = 16
7  first_activation = "relu"
8  final_activation = "sigmoid"
9  optimizer = "rmsprop"
10 loss = 'binary_crossentropy'
11 model = build_binary_output_model(
12     shape=model_shape,
13     layers=layers,
14     hidden_units = hidden_units,
15     first_activation = first_activation,
16     final_activation=final_activation,
17     optimizer=optimizer,
18     loss=loss, metrics=['accuracy']
19 )
20
21 model.fit(x_train, y_train, epochs=20, batch_size=512)
22 results = model.evaluate(x_test, y_test)
23
```

```
Epoch 1/20
49/49 [=====] - 2s 23ms/step - loss: 0.4584
- accuracy: 0.8221
Epoch 2/20
49/49 [=====] - 1s 11ms/step - loss: 0.2575
- accuracy: 0.9127
Epoch 3/20
49/49 [=====] - 1s 12ms/step - loss: 0.2012
- accuracy: 0.9289
Epoch 4/20
49/49 [=====] - 1s 12ms/step - loss: 0.1671
- accuracy: 0.9422
Epoch 5/20
49/49 [=====] - 1s 11ms/step - loss: 0.1470
- accuracy: 0.9476
Epoch 6/20
49/49 [=====] - 1s 11ms/step - loss: 0.1283
- accuracy: 0.9560
Epoch 7/20
49/49 [=====] - 1s 11ms/step - loss: 0.1127
- accuracy: 0.9611
Epoch 8/20
49/49 [=====] - 0s 10ms/step - loss: 0.1001
- accuracy: 0.9661
Epoch 9/20
49/49 [=====] - 0s 10ms/step - loss: 0.0855
- accuracy: 0.9720
Epoch 10/20
49/49 [=====] - 0s 9ms/step - loss: 0.0766 -
accuracy: 0.9756
Epoch 11/20
49/49 [=====] - 0s 10ms/step - loss: 0.0677
- accuracy: 0.9780
Epoch 12/20
49/49 [=====] - 0s 10ms/step - loss: 0.0588
- accuracy: 0.9821
Epoch 13/20
49/49 [=====] - 0s 9ms/step - loss: 0.0517 -
accuracy: 0.9849
Epoch 14/20
49/49 [=====] - 0s 8ms/step - loss: 0.0423 -
accuracy: 0.9883
```

```
Out[19]: [0.8005822896957397, 0.8503599762916565]
```

Further experiments

The following experiments will help convince you that the architecture choices you've made are all fairly reasonable, although there's still room for improvement:

- You used two hidden layers. Try using one or three hidden layers, and see how doing so affects validation and test accuracy.
- Try using layers with more hidden units or fewer hidden units: 32 units, 64 units, and so on.
- Try using the mse loss function instead of binary_crossentropy.

- Try using the tanh activation (an activation that was popular in the early days of neural networks) instead of relu.

```
In [20]: 1  #!/*****  
2  #!/*** Suggestion use One Hidden Layer  
3  #!/*****  
4  model_shape = (10000,)   
5  layers = 2  
6  hidden_units = 16  
7  first_activation = "relu"  
8  final_activation = "sigmoid"  
9  optimizer = "rmsprop"  
10 loss = 'binary_crossentropy'  
11 model = build_binary_output_model(  
12     shape=model_shape,  
13     layers=layers,  
14     hidden_units = hidden_units,  
15     first_activation = first_activation,  
16     final_activation=final_activation,  
17     optimizer=optimizer,  
18     loss=loss,metrics=['accuracy']  
19 )  
20  
21 model.fit(x_train, y_train, epochs=20, batch_size=512)  
22 results = model.evaluate(x_test, y_test)  
23
```

```

Epoch 1/20
49/49 [=====] - 1s 16ms/step - loss: 0.4495
- accuracy: 0.8324
Epoch 2/20
49/49 [=====] - 1s 16ms/step - loss: 0.2785
- accuracy: 0.9081
Epoch 3/20
49/49 [=====] - 0s 8ms/step - loss: 0.2204 -
accuracy: 0.9239
Epoch 4/20
49/49 [=====] - 0s 9ms/step - loss: 0.1874 -
accuracy: 0.9370
Epoch 5/20
49/49 [=====] - 0s 10ms/step - loss: 0.1650
- accuracy: 0.9436
Epoch 6/20
49/49 [=====] - 0s 9ms/step - loss: 0.1493 -
accuracy: 0.9503
Epoch 7/20
49/49 [=====] - 1s 10ms/step - loss: 0.1349
- accuracy: 0.9552
Epoch 8/20
49/49 [=====] - 0s 10ms/step - loss: 0.1243
- accuracy: 0.9589
Epoch 9/20
49/49 [=====] - 1s 13ms/step - loss: 0.1146
- accuracy: 0.9626
Epoch 10/20
49/49 [=====] - 0s 10ms/step - loss: 0.1049
- accuracy: 0.9666
Epoch 11/20
49/49 [=====] - 0s 10ms/step - loss: 0.0973
- accuracy: 0.9700

```

Out[20]: [0.5385316610336304, 0.8536800146102905]

```

In [21]: ▶ 1  #!/*****
2  #!/*** Suggestion: use Three Hidden Layers
3  #!/*****
4  model_shape = (10000,)
5  layers = 4
6  hidden_units = 16
7  first_activation = "relu"
8  final_activation = "sigmoid"
9  optimizer = "rmsprop"
10 loss = 'binary_crossentropy'
11 model = build_binary_output_model(
12     shape=model_shape,
13     layers=layers,
14     hidden_units = hidden_units,
15     first_activation = first_activation,
16     final_activation=final_activation,
17     optimizer=optimizer,
18     loss=loss,metrics=['accuracy']
19 )
20

```

```
21 model.fit(x_train, y_train, epochs=20, batch_size=512)
22 results = model.evaluate(x_test, y_test)
23
```

Epoch 1/20

49/49 [=====] - 5s 12ms/step - loss: 0.4572
- accuracy: 0.8080

Epoch 2/20

49/49 [=====] - 1s 11ms/step - loss: 0.2557
- accuracy: 0.9085

Epoch 3/20

49/49 [=====] - 0s 7ms/step - loss: 0.1973 -
accuracy: 0.9289

Epoch 4/20

49/49 [=====] - 0s 8ms/step - loss: 0.1648 -
accuracy: 0.9420

Epoch 5/20

49/49 [=====] - 0s 10ms/step - loss: 0.1415
- accuracy: 0.9494

Epoch 6/20

49/49 [=====] - 0s 10ms/step - loss: 0.1246
- accuracy: 0.9551

Epoch 7/20

49/49 [=====] - 0s 10ms/step - loss: 0.1048
- accuracy: 0.9638

Epoch 8/20

49/49 [=====] - 0s 9ms/step - loss: 0.0933 -
accuracy: 0.9686

Epoch 9/20

49/49 [=====] - 1s 10ms/step - loss: 0.0817
- accuracy: 0.9736 0s - loss: 0.0794 - accuracy: 0.

Epoch 10/20

49/49 [=====] - 0s 8ms/step - loss: 0.0671 -
accuracy: 0.9793

Epoch 11/20

49/49 [=====] - 1s 11ms/step - loss: 0.0574
- accuracy: 0.9816

Epoch 12/20

49/49 [=====] - 1s 12ms/step - loss: 0.0459
- accuracy: 0.9872

Epoch 13/20

49/49 [=====] - 0s 9ms/step - loss: 0.0402 -
accuracy: 0.9889

Epoch 14/20

49/49 [=====] - 0s 8ms/step - loss: 0.0352 -
accuracy: 0.9893

Epoch 15/20

49/49 [=====] - 0s 10ms/step - loss: 0.0277
- accuracy: 0.9921

Epoch 16/20

49/49 [=====] - 1s 11ms/step - loss: 0.0219
- accuracy: 0.9941

Epoch 17/20

49/49 [=====] - 0s 9ms/step - loss: 0.0162 -
accuracy: 0.9958

Epoch 18/20

49/49 [=====] - 1s 12ms/step - loss: 0.0163

```
- accuracy: 0.9956
Epoch 19/20
49/49 [=====] - 1s 12ms/step - loss: 0.0130
- accuracy: 0.9963
Epoch 20/20
49/49 [=====] - 0s 9ms/step - loss: 0.0105 -
accuracy: 0.9970
782/782 [=====] - 29s 38ms/step - loss: 0.88
51 - accuracy: 0.8500
```

Out[21]: [0.8851346969604492, 0.8499600291252136]

```
In [22]: ▶ 1  #!/*****
2  #!/*** Suggestion: use 32 Hidden Units
3  #!/*****
4  model_shape = (10000,)
5  layers = 3
6  hidden_units = 32
7  first_activation = "relu"
8  final_activation = "sigmoid"
9  optimizer = "rmsprop"
10 loss = 'binary_crossentropy'
11 model = build_binary_output_model(
12     shape=model_shape,
13     layers=layers,
14     hidden_units = hidden_units,
15     first_activation = first_activation,
16     final_activation=final_activation,
17     optimizer=optimizer,
18     loss=loss,metrics=['accuracy']
19 )
20
21 model.fit(x_train, y_train, epochs=20, batch_size=512)
22 results = model.evaluate(x_test, y_test)
23
```

```

Epoch 1/20
49/49 [=====] - 1s 13ms/step - loss: 0.4266
- accuracy: 0.8257
Epoch 2/20
49/49 [=====] - 1s 11ms/step - loss: 0.2386
- accuracy: 0.9136
Epoch 3/20
49/49 [=====] - 1s 14ms/step - loss: 0.1878
- accuracy: 0.9325
Epoch 4/20
49/49 [=====] - 1s 13ms/step - loss: 0.1589
- accuracy: 0.9416
Epoch 5/20
49/49 [=====] - 1s 13ms/step - loss: 0.1383
- accuracy: 0.9494
Epoch 6/20
49/49 [=====] - 1s 12ms/step - loss: 0.1209
- accuracy: 0.9580
Epoch 7/20
49/49 [=====] - 1s 12ms/step - loss: 0.1022
- accuracy: 0.9631
Epoch 8/20
49/49 [=====] - 1s 14ms/step - loss: 0.0927
- accuracy: 0.9674
Epoch 9/20
49/49 [=====] - 1s 16ms/step - loss: 0.0812
- accuracy: 0.9720 0s - loss: 0.0737 - accuracy
Epoch 10/20
49/49 [=====] - 1s 14ms/step - loss: 0.0681
- accuracy: 0.9779
Epoch 11/20
49/49 [=====] - 1s 13ms/step - loss: 0.0606
- accuracy: 0.9811
Epoch 12/20
49/49 [=====] - 1s 12ms/step - loss: 0.0499
- accuracy: 0.9839
Epoch 13/20
49/49 [=====] - 1s 11ms/step - loss: 0.0393
- accuracy: 0.9883
Epoch 14/20
49/49 [=====] - 1s 13ms/step - loss: 0.0352
- accuracy: 0.9893

```

Out[22]: [0.9207907319068909, 0.8514400124549866]

In [23]:

```

1  #//*****
2  #//*** Suggestion: use 64 Hidden Units
3  #//*****
4  model_shape = (10000,)
5  layers = 3
6  hidden_units = 64
7  first_activation = "relu"
8  final_activation = "sigmoid"
9  optimizer = "rmsprop"
10 loss = 'binary_crossentropy'
11 model = build_binary_output_model(
12     shape=model_shape,

```



```
13     layers=layers,
14     hidden_units = hidden_units,
15     first_activation = first_activation,
16     final_activation=final_activation,
17     optimizer=optimizer,
18     loss=loss,metrics=['accuracy']
19 )
20
21 model.fit(x_train, y_train, epochs=5, batch_size=512)
22 results = model.evaluate(x_test, y_test)
23
```

Epoch 1/5

49/49 [=====] - 1s 16ms/step - loss: 0.4240
- accuracy: 0.8107

Epoch 2/5

49/49 [=====] - 1s 18ms/step - loss: 0.2395
- accuracy: 0.9076

Epoch 3/5

49/49 [=====] - 1s 16ms/step - loss: 0.1847
- accuracy: 0.9298

Epoch 4/5

49/49 [=====] - 1s 20ms/step - loss: 0.1458
- accuracy: 0.9434

Epoch 5/5

49/49 [=====] - 1s 19ms/step - loss: 0.1128
- accuracy: 0.9589
782/782 [=====] - 31s 40ms/step - loss: 0.3607
- accuracy: 0.8748

Out[23]: [0.36067381501197815, 0.8748000264167786]

In [24]:

```
1  #!/*****
2  #!/*** Suggestion: use tanh activation
3  #!/*****
4  model_shape = (10000,)
5  layers = 3
6  hidden_units = 16
7  first_activation = "tanh"
8  final_activation = "sigmoid"
9  optimizer = "rmsprop"
10 loss = 'binary_crossentropy'
11 model = build_binary_output_model(
12     shape=model_shape,
13     layers=layers,
14     hidden_units = hidden_units,
15     first_activation = first_activation,
16     final_activation=final_activation,
17     optimizer=optimizer,
18     loss=loss,metrics=['accuracy']
19 )
20
21 model.fit(x_train, y_train, epochs=5, batch_size=512)
22 results = model.evaluate(x_test, y_test)
23
24 results
```

```
Epoch 1/5
49/49 [=====] - 4s 11ms/step - loss: 0.4162
- accuracy: 0.8337
Epoch 2/5
49/49 [=====] - 0s 10ms/step - loss: 0.2328
- accuracy: 0.9142
Epoch 3/5
49/49 [=====] - 1s 11ms/step - loss: 0.1780
- accuracy: 0.9352
Epoch 4/5
49/49 [=====] - 0s 9ms/step - loss: 0.1474 -
accuracy: 0.9478
Epoch 5/5
49/49 [=====] - 0s 8ms/step - loss: 0.1265 -
accuracy: 0.9558
782/782 [=====] - 29s 37ms/step - loss: 0.37
```

Out[24]: [0.3727782070636749, 0.867680013179779]

In [25]: ▶

```
1  #//*****
2  #//*** Suggestion: use MSE loss
3  #//*****
4  model_shape = (10000,)
5  layers = 3
6  hidden_units = 16
7  first_activation = "relu"
8  final_activation = "sigmoid"
9  optimizer = "rmsprop"
10 loss = 'mse'
11 model = build_binary_output_model(
12     shape=model_shape,
13     layers=layers,
14     hidden_units = hidden_units,
15     first_activation = first_activation,
16     final_activation=final_activation,
17     optimizer=optimizer,
18     loss=loss,metrics=['accuracy']
19 )
20
21 model.fit(x_train, y_train, epochs=5, batch_size=512)
22 results = model.evaluate(x_test, y_test)
23
```

```
Epoch 1/5
49/49 [=====] - 1s 10ms/step - loss: 0.1479
- accuracy: 0.8266
Epoch 2/5
```

Out[25]: [0.0892355889081955, 0.8794400095939636]

In []: 