

Stoneburner, Kurt

. DSC 650 - Assignment 5.3 - Tensorflow Keras Regression Example

https://github.com/fchollet/deep-learning-with-python-notebooks/blob/master/chapter04_getting-started-with-neural-networks.ipynb (https://github.com/fchollet/deep-learning-with-python-notebooks/blob/master/chapter04_getting-started-with-neural-networks.ipynb)

```
In [1]: ▶ 1 import os
2 import sys
3 # Imports and Load Data
4 import matplotlib.pyplot as plt
5 import numpy as np
6 import pandas as pd
7
8 Use the whole window in the IPYNB editor
9 from IPython.core.display import display, HTML
10 display(HTML("<style>.container { width:100% !important; }</style>"))
11
12 Maximize columns and rows displayed by pandas
13 pd.set_option('display.max_rows', 100)
14 pd.set_option('display.max_columns', None)
15
16
17 from tensorflow import keras
```

```
In [2]: ▶ 1 *****
2 Attempt to modularize a Sequential Keras Module.
3 *****
4 def build_model(**kwargs):
5     Define the Model
6     from tensorflow.keras import models
7     from tensorflow.keras import layers
8     from tensorflow.keras import optimizers
9
10
11     *****
12     Set Default values
13     *****
14
15     Total Layers is the total layers including the input 1
16     total_layers = 2
17
18     Hidden units to be applied to all layers except the las
19     hidden_units = 16
20     Activation Function to be applied to all layers except
21     first_activation = "relu"
22
```

```

23     #!/*** Activation Function for the last layer. No Activation
24     final_activation=None
25
26     #!/*** Complie Optimzer
27     optimizer='rmsprop'
28
29     #!/*** Loss Function for Optimizer
30     loss = 'mse'
31
32     #!/*** Optimizer Metrics
33     metrics=['accuracy']
34
35     #!/*** Tuple Defined Shape of the First Layer. None means thi
36     shape = None
37
38     #!/*** Apply compiler to the model
39     do_compile = True
40
41     #!/*** Number of Units (outputs) for the output layer.
42     output_layer = 1
43
44     #!/*** Print a Representation of the Model
45     display_model = True
46
47     #!/*** Modify the default settings with **kwargs
48     #!/*** Apply Kwargs
49     for key,value in kwargs.items():
50
51         if key == 'layers':
52             total_layers=value
53
54         if key == 'hidden_units':
55             hidden_units=value
56
57         if key == 'loss':
58             loss=value
59
60         if key == 'first_activation':
61             first_activation=value
62
63         if key == 'final_activation':
64             final_activation=value
65
66         if key == 'optimizer':
67             optimizer=value
68
69         if key == 'metrics':
70             metrics=value
71
72         if key == 'shape':
73             shape = value
74
75         if key == 'compile':
76             do_compile = value
77
78         if key == 'output_layer':

```

```

79         output_layer = value
80
81     if key == 'display_model':
82         display_model = value
83
84     disp = ""
85     #!/*** Initialize the model
86     model = models.Sequential()
87
88     if display_model:
89         disp += "models.Sequential()"
90         disp += "\n"
91
92     #!/*** Add the First Layer. Include an Input_Shape paramter i
93     if shape == None:
94         #!/*** Add First Layer
95         model.add(layers.Dense(hidden_units, activation=first_act
96         if display_model:
97             disp += f"model.add(layers.Dense({hidden_units}, acti
98             disp += "\n"
99
100     else:
101         #!/*** Add First Layer
102         model.add(layers.Dense(hidden_units, activation=first_act
103
104         if display_model:
105             disp += f"model.add(layers.Dense({hidden_units}, acti
106             disp += "\n"
107
108
109     #!/*** Add Additional Layers if total_layers greater than 2
110     for x in range(total_layers-2):
111
112         #!/*** These are basic layers with same number of hidden
113         model.add(layers.Dense(hidden_units, activation=first_act
114         if display_model:
115             disp += f"model.add(layers.Dense({hidden_units}, acti
116             disp += "\n"
117
118
119
120     #!/*** Add Final Layer
121
122     if final_activation == None:
123         model.add(layers.Dense(output_layer))
124
125         if display_model:
126             disp += f"model.add(layers.Dense({output_layer}))"
127             disp += "\n"
128
129     else:
130         model.add(layers.Dense(output_layer, activation=final_act
131
132         if display_model:
133             disp += f"model.add(layers.Dense({output_layer}, acti
134             disp += "\n"

```

```

135
136     #!/*** Compile Model
137     if do_compile:
138
139         model.compile(optimizer=optimizer,loss=loss,metrics=metri
140
141         if display_model:
142             disp += f"model.compile(optimizer={optimizer},loss={l
143             disp += "\n"
144
145         if display_model:
146             print(disp)
147
148     return model
149

```

Predicting house prices: A regression example

```

In [3]: 1 #!/*** Import the housing Dataset from Keras
        2 from tensorflow.keras.datasets import boston_housing

```

Type *Markdown* and LaTeX: α^2

```

In [4]: 1 print("Train Shape: ", train_data.shape)

```

```

Train Shape: (404, 13)
Test Shape: (102, 13)

```

Need Notes on normalizing the Data

And each feature in the input data (for example, the crime rate) has a different scale. For instance, some values are proportions, which take values between 0 and 1; others take values between 1 and 12, others between 0 and 100, and so on.

As you can see, you have 404 training samples and 102 test samples, each with 13 numerical features, such as per capita crime rate, average number of rooms per dwelling, accessibility to highways, and so on. The targets are the median values of owner-occupied homes, in thousands of dollars:

Normalization

feature-wise normalization: Subtract the Feature (column) Mean and Divide by the Standard Deviation

It would be problematic to feed into a neural network values that all take wildly different ranges. The network might be able to automatically adapt to such heterogeneous data, but it would definitely make learning more difficult. A widespread best practice to deal with such data is to do feature-wise normalization: for each feature in the input data (a column in the input data matrix), you subtract the mean of the feature and divide by the standard deviation, so that the feature is centered around 0 and has a unit standard deviation. This is easily done in Numpy.

Note: that the quantities used for normalizing the test data are computed using the training data.

You should never use in your workflow any quantity computed on the test data, even for something as simple as data normalization.

```
In [5]: 1 #!/*****
2 #!/*** Normalize the Data
3 #!/*****
4
5 #!/*** Get the Mean of each Column in the Matrix
6 mean = train_data.mean(axis=0)
7
8 #!/*** Subtract the Mean
9 train_data -= mean
10
11 #!/*** Get the Standard Deviation
12 std = train_data.std(axis=0)
13
14 #!/*** Divide by the Standard Deviation
15 train_data /= std
16
17 #!/*** Subtract the Mean from the test Data
18 test_data -= mean
19
20 #!/*** Divide the Standard Deviation from the Test Data
```

Because so few samples are available, you'll use a very small network with two hidden layers, each with 64 units. In general, the less training data you have, the worse overfitting will be, and using a small network is one way to mitigate overfitting.

The network ends with a single unit and no activation (it will be a linear layer). This is a typical setup for scalar regression (a regression where you're trying to predict a single continuous value). Applying an activation function would constrain the range the out-put can take; for instance, if you applied a sigmoid activation function to the last layer, the network could only learn to predict values between 0 and 1. Here, because the last layer is purely linear, the network is free to learn to predict values in any range.

Note: that you compile the network with the mse loss function—mean squared error, the square of the difference between the predictions and the targets. This is a widely used loss function for regression problems. You're also monitoring a new metric during training: mean absolute error (MAE). It's the absolute value of the difference between the predictions and the targets. For instance, an MAE of 0.5 on this problem would mean your predictions are off by \$500 on average.

Validate the Model ysubf K-fold validation

To evaluate your network while you keep adjusting its parameters (such as the number of epochs used for training), you could split the data into a training set and a validation set, as you did in the previous examples. But because you have so few data points, the validation set would end up being very small (for instance, about 100 examples). As a consequence, the validation scores might change a lot depending on which datapoints you chose to use for validation and which you chose for training: the validation scores might have a high variance with regard to the validation split. This would prevent you from reliably evaluating your model. The best practice in such

situations is to use K-fold cross-validation. It consists of splitting the available data into K partitions (typically $K = 4$ or 5), instantiating K identical models, and training each one on $K - 1$ partitions while evaluating on the remaining partition. The validation score for the model used is then the average of the K validation scores obtained. In terms of code, this is straightforward.

Instead of selecting a Validation Subset, K-Fold validation runs splits the data into multiple different validation sets and scores them all. The average of the validations is a validation estimate of the model.

```
In [6]: 1  #//*****
2  #//*** Book Supplied Model Settings
3  #//*****
4  layers = 3
5  hidden_units = 64
6  first_activation = "relu"
7  optimizer = "rmsprop"
8  loss = 'mse'
9  metrics = ['mae']
10
11
12  #//*** F-Fold Validation
13  k = 4
14  num_val_samples = len(train_data) // k #//*** Floor Division
15  num_epochs = 100
16  all_scores = []
17  for i in range(k):
18      print(f"Processing fold #{i}")
19      val_data = train_data[i * num_val_samples: (i + 1) * num_val_s
20      val_targets = train_targets[i * num_val_samples: (i + 1) * num
21      partial_train_data = np.concatenate(
22          [train_data[:i * num_val_samples],
23           train_data[(i + 1) * num_val_samples:]],
24          axis=0)
25      partial_train_targets = np.concatenate(
26          [train_targets[:i * num_val_samples],
27           train_targets[(i + 1) * num_val_samples:]],
28          axis=0)
29
30      model = build_model(
31          layers=layers,
32          hidden_units = hidden_units,
33          first_activation = first_activation,
34          optimizer=optimizer,
35          loss=loss,
36          metrics=metrics,
37          output_layer = 1
38      )
39
40      model.fit(partial_train_data, partial_train_targets,
41                epochs=num_epochs, batch_size=16, verbose=0)
42      val_mse, val_mae = model.evaluate(val_data, val_targets, verbo
43      all_scores.append(val_mae)
44
45
```

```

Processing fold #0
models.Sequential()
model.add(layers.Dense(64, activation=relu ))
model.add(layers.Dense(64, activation=relu))
model.add(layers.Dense(1))
model.compile(optimizer=rmsprop,loss=mse,metrics=['mae'])

```

```

Processing fold #1
models.Sequential()
model.add(layers.Dense(64, activation=relu ))
model.add(layers.Dense(64, activation=relu))
model.add(layers.Dense(1))
model.compile(optimizer=rmsprop,loss=mse,metrics=['mae'])

```

```

Processing fold #2
models.Sequential()
model.add(layers.Dense(64, activation=relu ))
model.add(layers.Dense(64, activation=relu))
model.add(layers.Dense(1))
model.compile(optimizer=rmsprop,loss=mse,metrics=['mae'])

```

```

Processing fold #3
models.Sequential()
model.add(layers.Dense(64, activation=relu ))
model.add(layers.Dense(64, activation=relu))
model.add(layers.Dense(1))

```

In [7]: 1 print("All Scores: ", all_scores)

```

All Scores: [1.8729946613311768, 2.6107888221740723, 2.4159820079803
467, 2.37795090675354]
Mean Scores: 2.319429099559784

```

In [8]:

```

1  #!/*** Saving the validation logs of each fold
2  num_epochs = 500
3  all_mae_histories = []
4  for i in range(k):
5      print(f"Processing fold #{i}")
6      val_data = train_data[i * num_val_samples: (i + 1) * num_val_s
7      val_targets = train_targets[i * num_val_samples: (i + 1) * num
8      partial_train_data = np.concatenate(
9          [train_data[:i * num_val_samples],
10           train_data[(i + 1) * num_val_samples:]],
11          axis=0)
12      partial_train_targets = np.concatenate(
13          [train_targets[:i * num_val_samples],
14           train_targets[(i + 1) * num_val_samples:]],
15          axis=0)
16      model = build_model(
17          layers=layers,
18          hidden_units = hidden_units,
19          first_activation = first_activation,
20          optimizer=optimizer,
21          loss=loss,
22          metrics=metrics,

```

```

23         output_layer = 1
24     )
25     history = model.fit(partial_train_data, partial_train_targets,
26                         validation_data=(val_data, val_targets),
27                         epochs=num_epochs, batch_size=16, verbose=
28
29     mae_history = history.history["val_mae"]
30     all_mae_histories.append(mae_history)
31

```

```

Processing fold #0
models.Sequential()
model.add(layers.Dense(64, activation=relu ))
model.add(layers.Dense(64, activation=relu))
model.add(layers.Dense(1))
model.compile(optimizer=rmsprop, loss=mse, metrics=['mae'])

```

```

Processing fold #1
models.Sequential()
model.add(layers.Dense(64, activation=relu ))
model.add(layers.Dense(64, activation=relu))
model.add(layers.Dense(1))
model.compile(optimizer=rmsprop, loss=mse, metrics=['mae'])

```

```

Processing fold #2
models.Sequential()
model.add(layers.Dense(64, activation=relu ))
model.add(layers.Dense(64, activation=relu))
model.add(layers.Dense(1))
model.compile(optimizer=rmsprop, loss=mse, metrics=['mae'])

```

```

Processing fold #3
models.Sequential()
model.add(layers.Dense(64, activation=relu ))
model.add(layers.Dense(64, activation=relu))
model.add(layers.Dense(1))
model.compile(optimizer=rmsprop, loss=mse, metrics=['mae'])

```

```



In [9]: 1  #!/*** Building the history of successive mean K-fold validation s
        2  average_mae_history = [

```

```

In [10]: 1  #!/*** Plotting Validation scores
        2  plt.plot(range(1, len(average_mae_history) + 1), average_mae_histo
        3  plt.xlabel("Epochs")
        4  plt.ylabel("Validation MAE")
        5  plt.show()
        6
        7  #!/*** Plotting Validation scores excluding the first 10 data Poin
        8
        9
       10  truncated_mae_history = average_mae_history[10:]
       11  plt.plot(range(1, len(truncated_mae_history) + 1), truncated_mae_h
       12  plt.xlabel("Epochs")
       13  plt.ylabel("Validation MAE")
       14  plt.show()

```

```
In [11]: ▶ 1 #!/** Training the Final Model  
2 #model = build_model()
```

```

3
4
5 #!/*****
6 #!/*** Book Supplied Settings
7 #!/*****
8 layers = 3
9 hidden_units = 64
10 first_activation = "relu"
11 optimizer = "rmsprop"
12 loss = 'mse'
13 metrics = ['mae']
14 model = build_model(
15     layers=layers,
16     hidden_units = hidden_units,
17     first_activation = first_activation,
18     optimizer=optimizer,
19     loss=loss,
20     metrics=metrics,
21     output_layer = 1
22 )
23
24 model.fit(train_data, train_targets,
25           epochs=130, batch_size=16, verbose=0)
26 test_mse_score, test_mae_score = model.evaluate(test_data, test_ta
27
28 print("test_mae_score: ",test_mae_score)
29
30 #!/*** Generate Predictions on the test_data
31 predictions = model.predict(test_data)
32 print("Predictions: ", predictions[0])
33
34 predictions = predictions.flatten()
35

```

```

models.Sequential()
model.add(layers.Dense(64, activation=relu ))
model.add(layers.Dense(64, activation=relu))
model.add(layers.Dense(1))
model.compile(optimizer=rmsprop,loss=mse,metrics=['mae'])

4/4 [=====] - 0s 0s/step - loss: 16.9921 - m
ae: 2.5225
test_mae_score: 2.5225348472595215
Predictions: [9.462795]

```

In []: 

```
In [15]: ▶ 1 #!/*** Plot Predicted vs actual Values
2 plt.plot(range(1, len(predictions) + 1), test_targets, "bo", label=
3 plt.plot(range(1, len(predictions) + 1), predictions, "b", label="
4 plt.ylabel("Prices")
5 plt.legend()
6 plt.title("Predicted vs Actual Housing Prices")
```

